Deep Learning and Optimization

Unpacking Transformers, LLMs and Diffusion

Session 1

olivier.koch@ensae.fr

slack #ensae-dl-2025

What you will learn

- Concepts and practical aspects of deep learning
- Build your own backprop, mini-gpt and diffusion models

This course is <u>not</u> a complete overview of all ML & deep learning techniques.

Learning from scratch matters

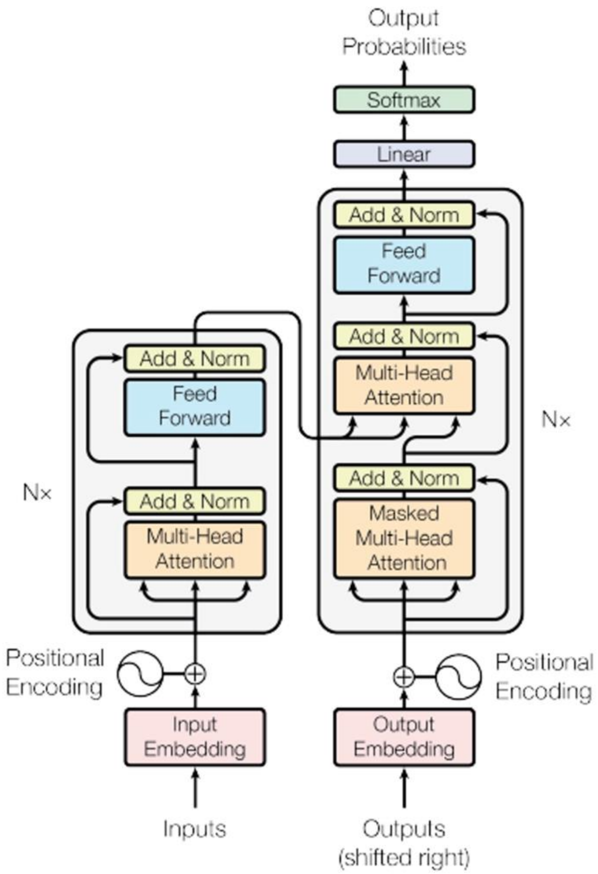

mmm… so what?

Practical stuff

- 6 sessions
- 1.5h theory + 1.5h practice
- Grading: 6 notebooks (66%) + 1 quiz (33%)
- I expect you to return your notebook <u>at the end of each session</u>.
- You can then send an updated version over the following week to aim for a better grade (and learn more).

Rules of engagement

- It’s OK to ask for help and share tips (use slack)
- Each notebook should be yours
- You must return your notebook as-is at the end of each TP (via slack)
- You can then send an updated version within a week

References

- Understanding Deep Learning, Simon Prince, December 2023
- Deep Learning: Foundations and Concepts, Christopher and Hugh Bishop, 2023
- Neural networks: Zero to Hero, Andrej Karpathy, Youtube Lecture Series, 2022

Session	Date	Topic
1	05-02-2025	Intro to Deep Learning Practical: micrograd
2	12-02-2025	DL fundamentals ● Backprop ● Loss functions Practical: bigram, MLP for next character prediction
3	19-02-2025	DL Fundamentals II ● Activation functions ● Regularization ● Initialization ● Residual networks ● Normalization Practical: tensor-based models
	26-02-2025	Pas de cours
4	05-03-2025	Attention & Transformers Practical: GPT from scratch
5	12-03-2025	DL for computer vision Practical: Convnets for CIFAR-10
6	19-03-2025	VAE & Diffusion models Practical: diffusion from scratch Quiz/Exam

Session	Date	Topic
1	05-02-2025	Intro to Deep Learning Practical: micrograd
2	12-02-2025	DL fundamentals • Backprop • Loss functions Practical: bigram, MLP for next character prediction
3	19-02-2025	DL Fundamentals II • Activation functions • Regularization • Initialization • Residual networks • Normalization Practical: tensor-based models
	26-02-2025	Pas de cours
4	05-03-2025	Attention & Transformers Practical: GPT from scratch
5	12-03-2025	DL for computer vision Practical: Convnets for CIFAR-10
6	19-03-2025	VAE & Diffusion models Practical: diffusion from scratch Quiz/Exam

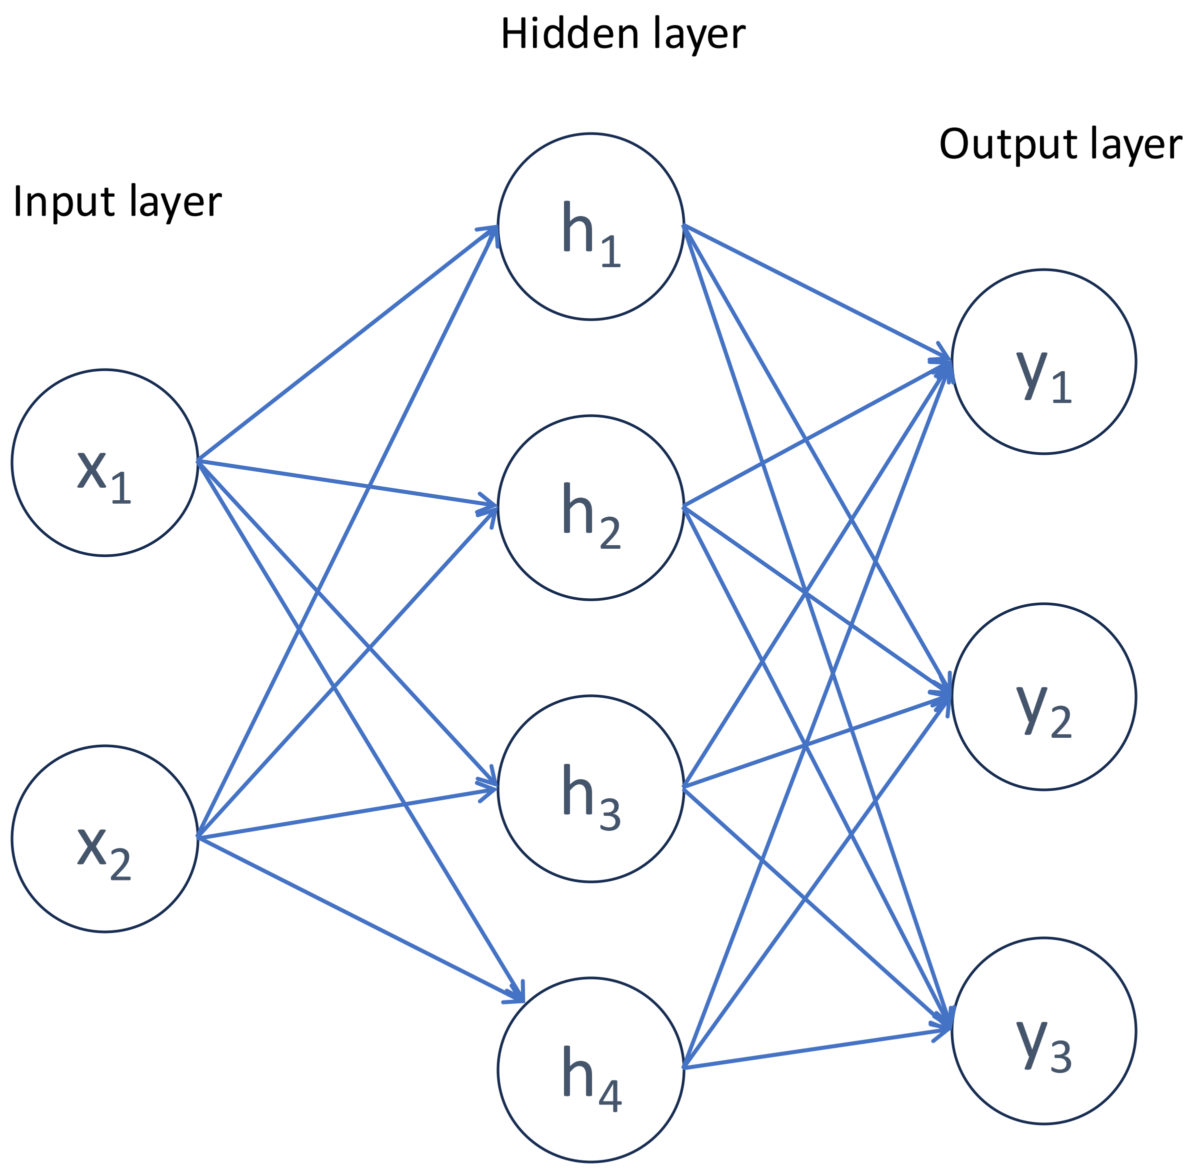
Hidden layer
Output layer
Input layer
h1
h2
h3
h4
x1
x2
y1
y2
y3

Linear transformation and non-linear activation

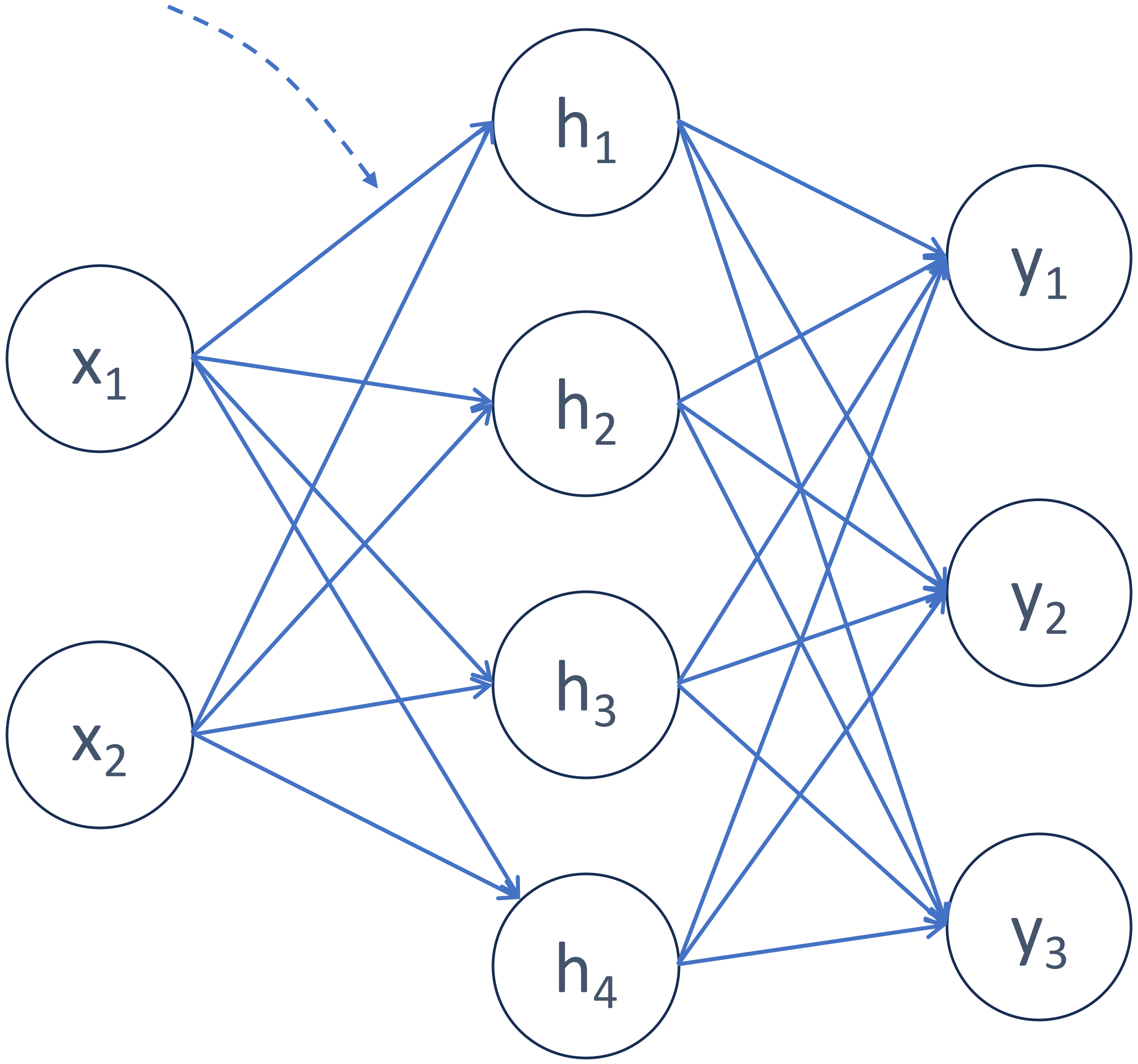

$$h_i = a(\beta_i + \sum \theta_{ij} x_j)$$

$$y_i = \gamma_i + \sum \varphi_{ij} h_j$$

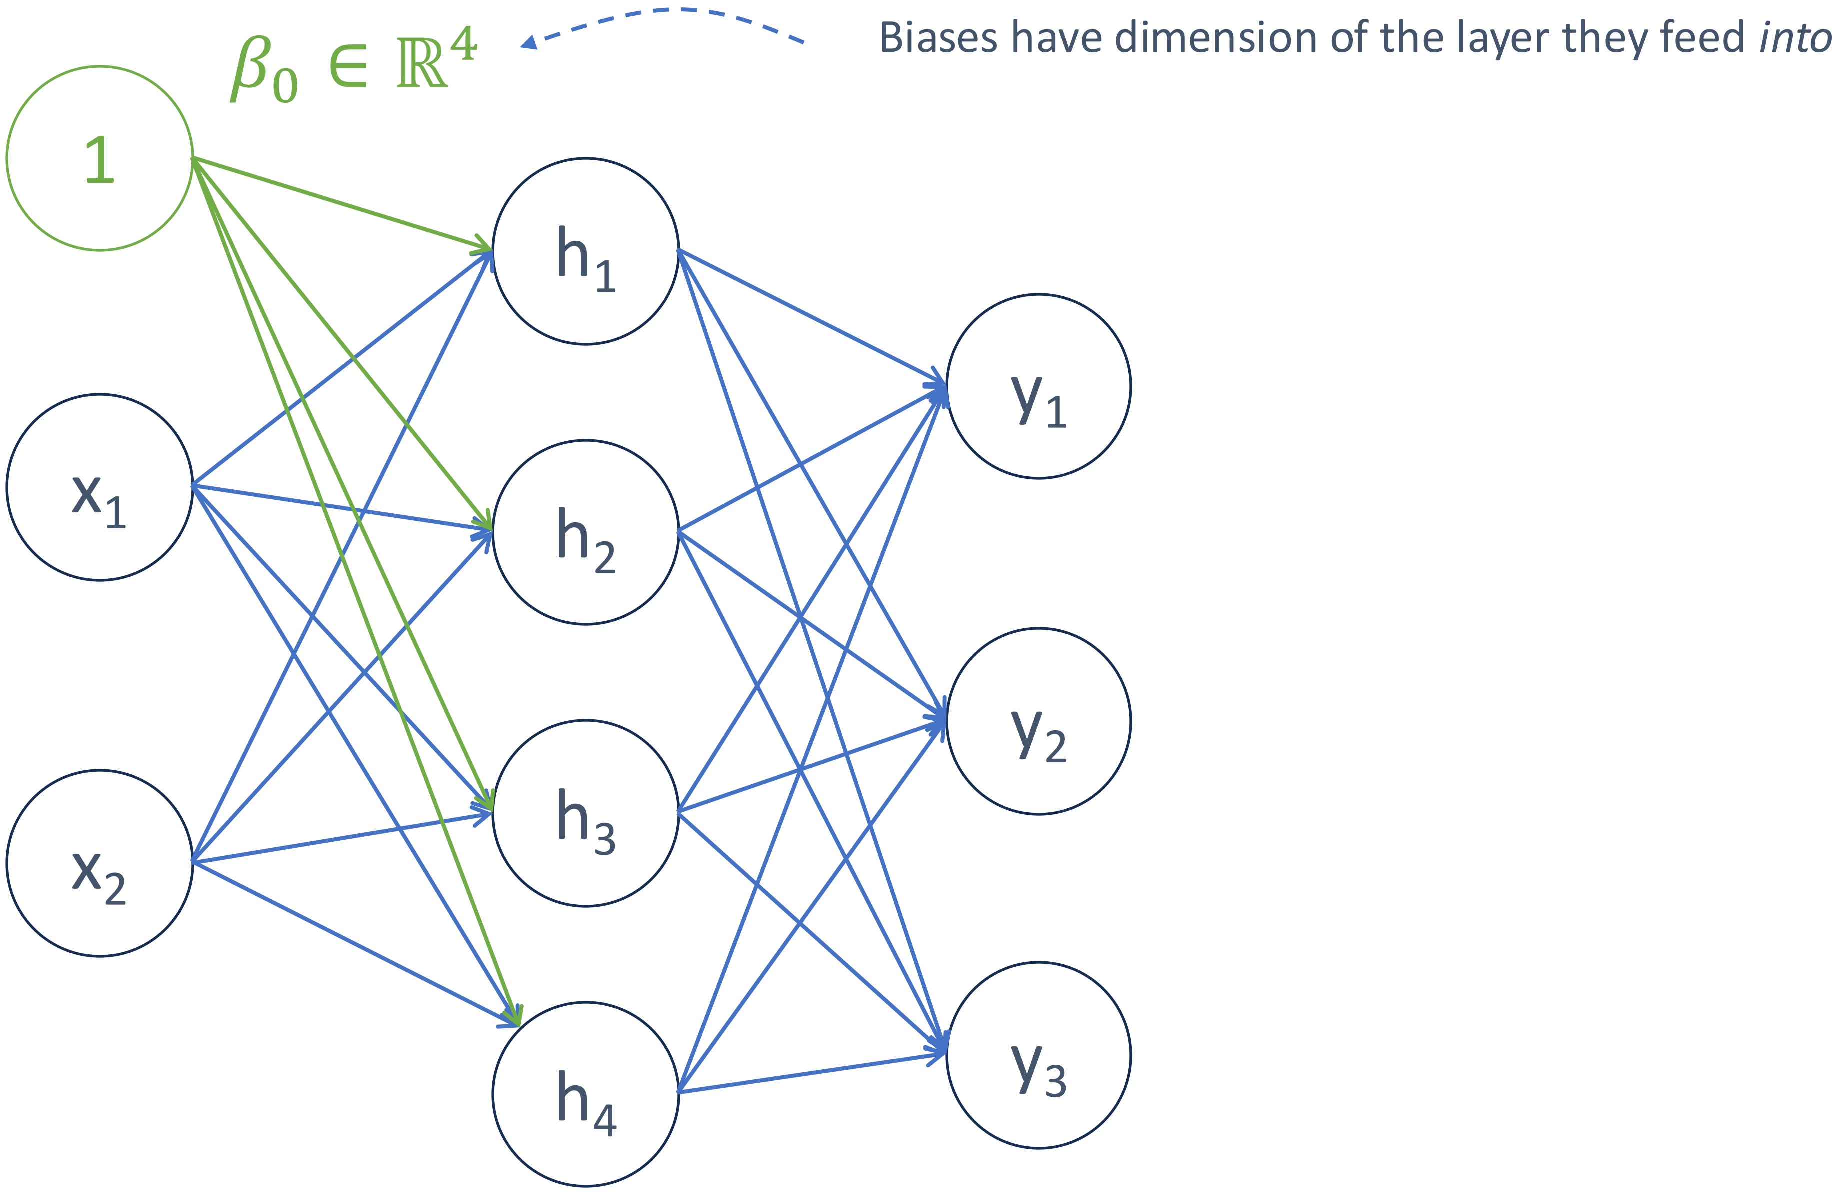

$\beta_0 \in \mathbb{R}^4$
Biases have dimension of the layer they feed *into*
1
x_1
x_2
h_1
h_2
h_3
h_4
y_1
y_2
y_3

Activation functions add non-linearity to the network. Will be discussed in Session 3!

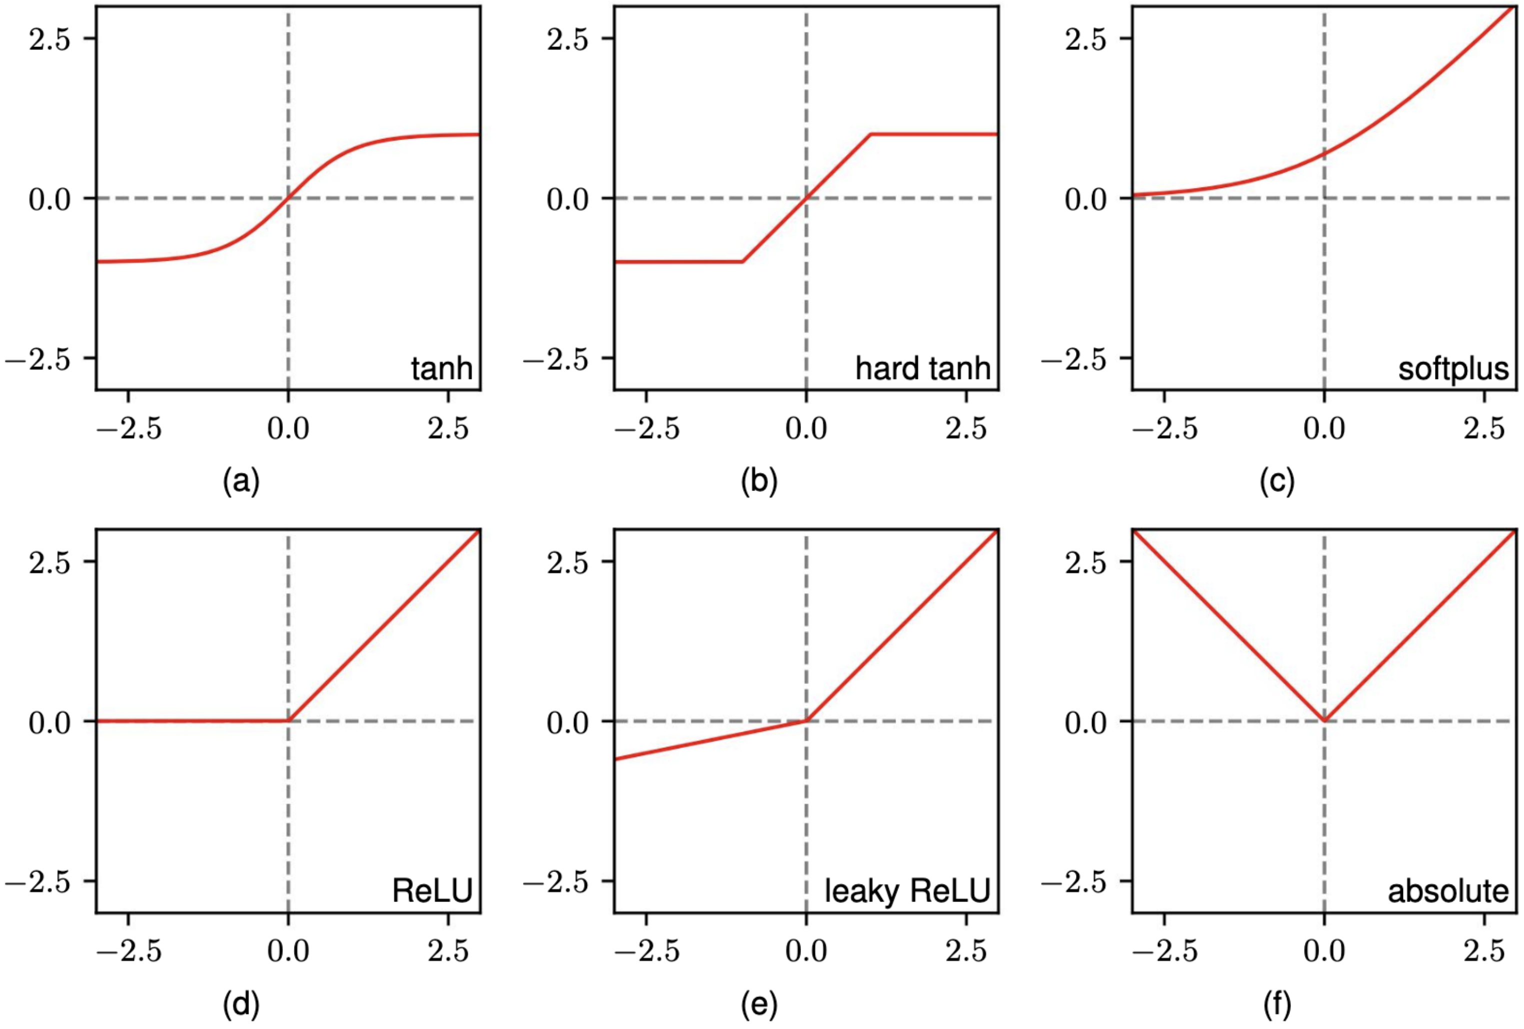

Figure 6.12 A variety of nonlinear activation functions.

Source: Deep Learning Foundations and Concept, Bishop, 2023

Backprop and gradient descent

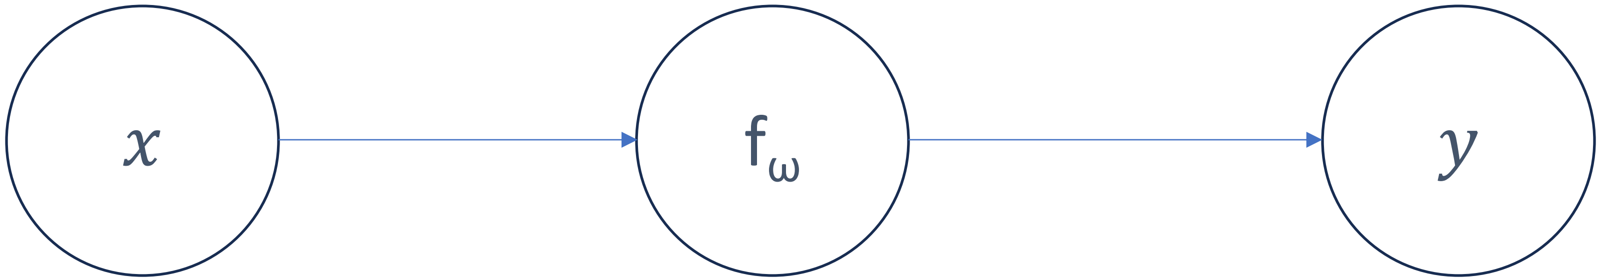

$$\ell = (y - \hat{y})^2$$

ground-truth

Backprop and gradient descent

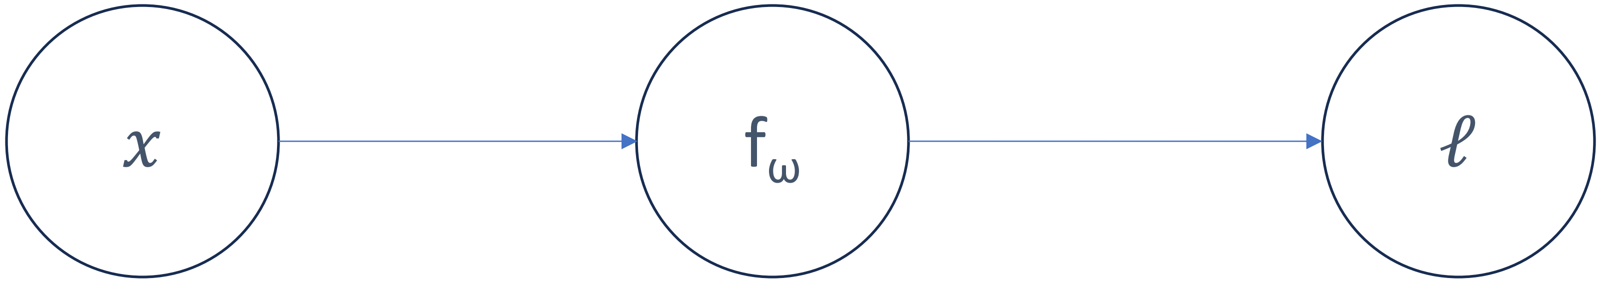

Goal: find ω that minimizes:

$$\ell = f_\omega(x)$$

Gradient descent

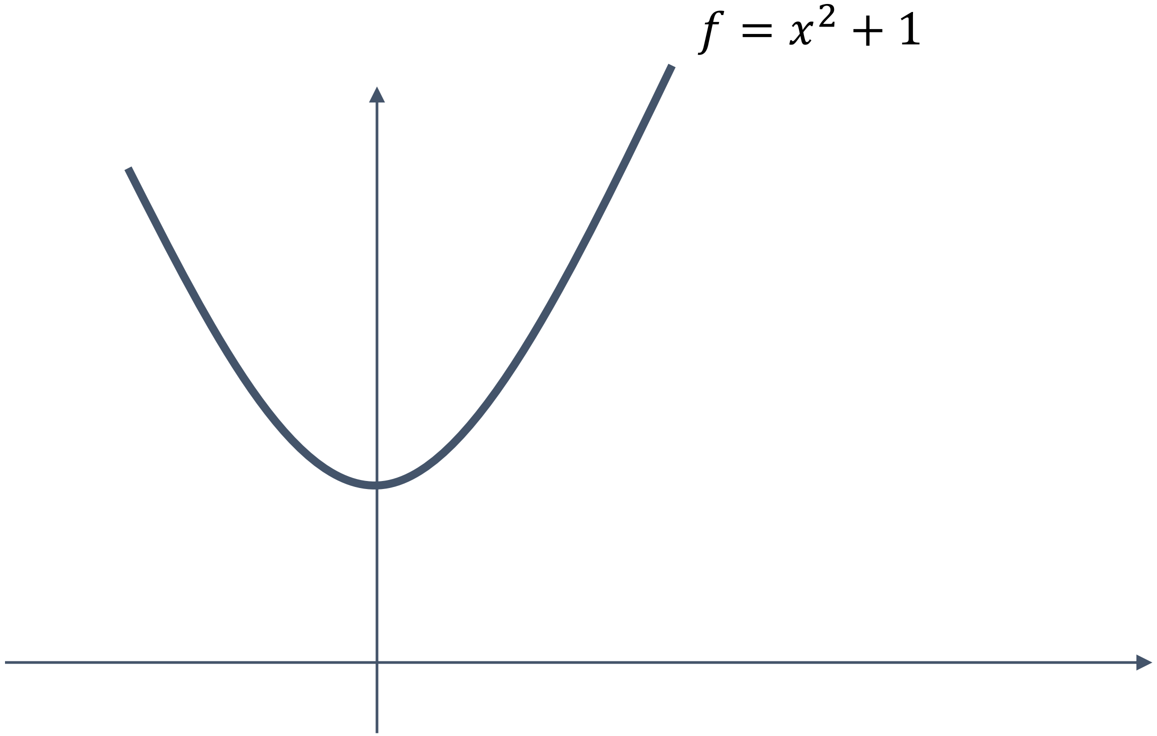

Gradient descent

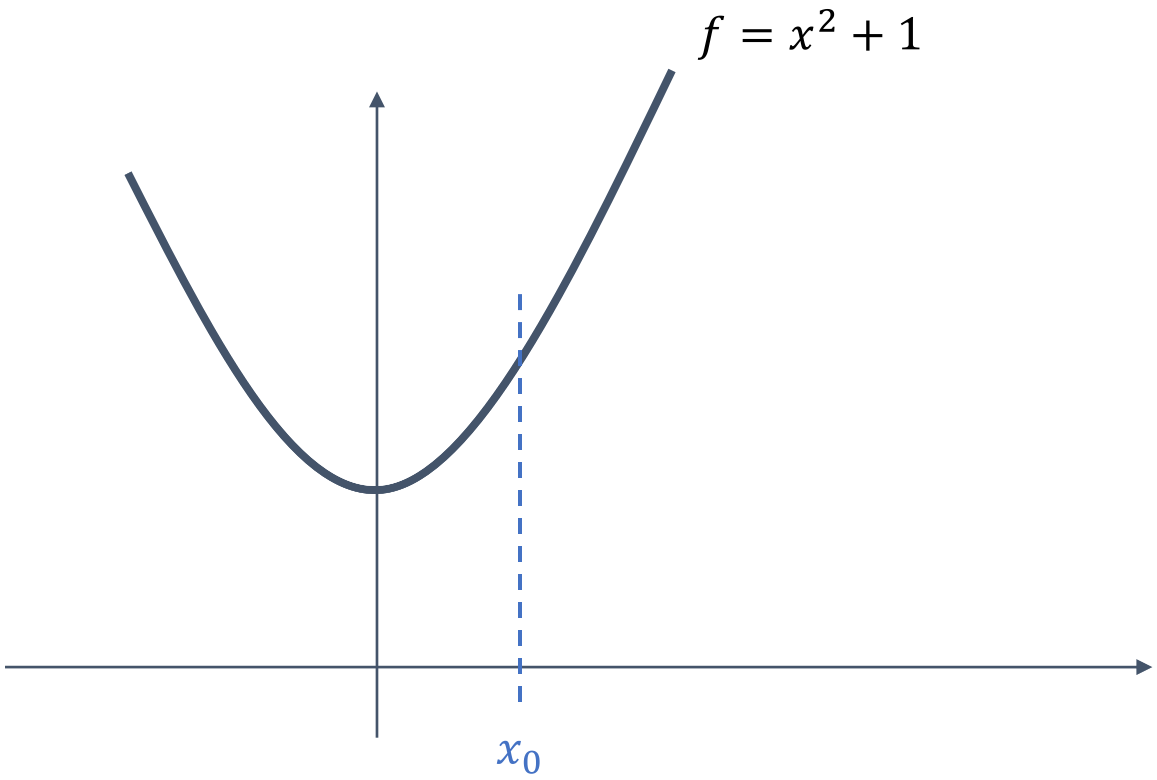

Gradient descent

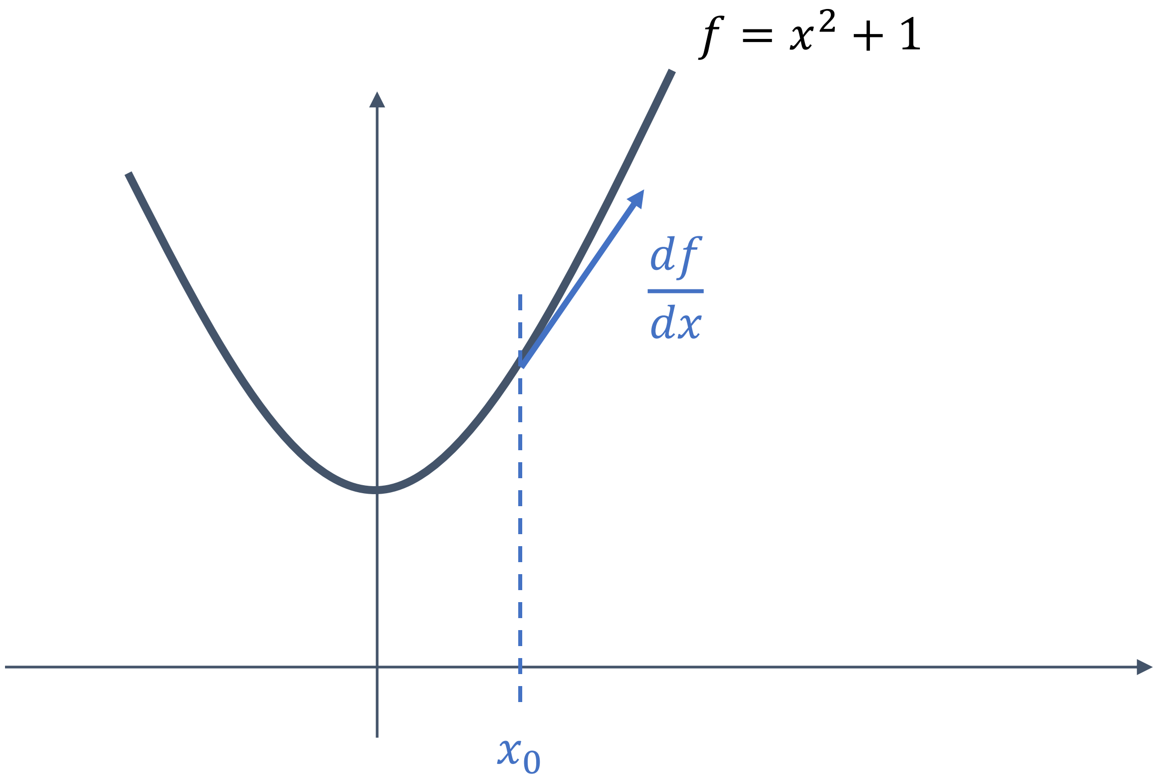

Gradient descent

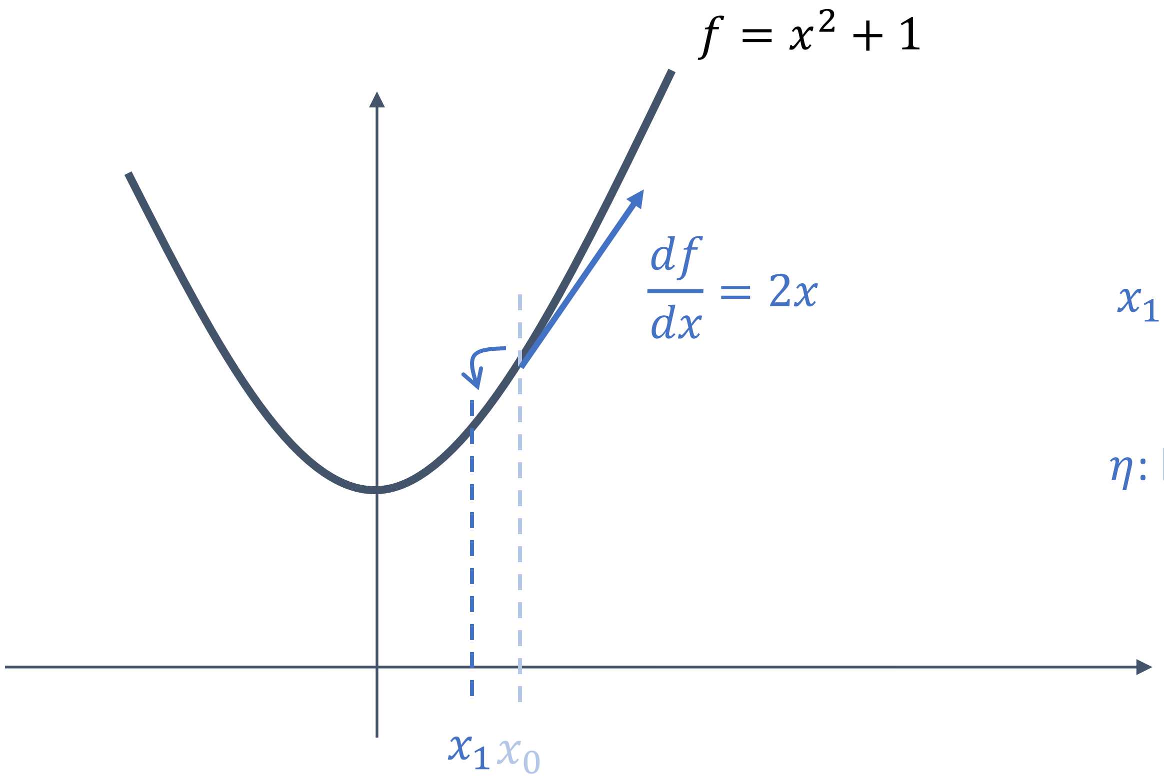

$$x_1 = x_0 - \eta \frac{df}{dx}$$

η: learning rate

Gradient descent

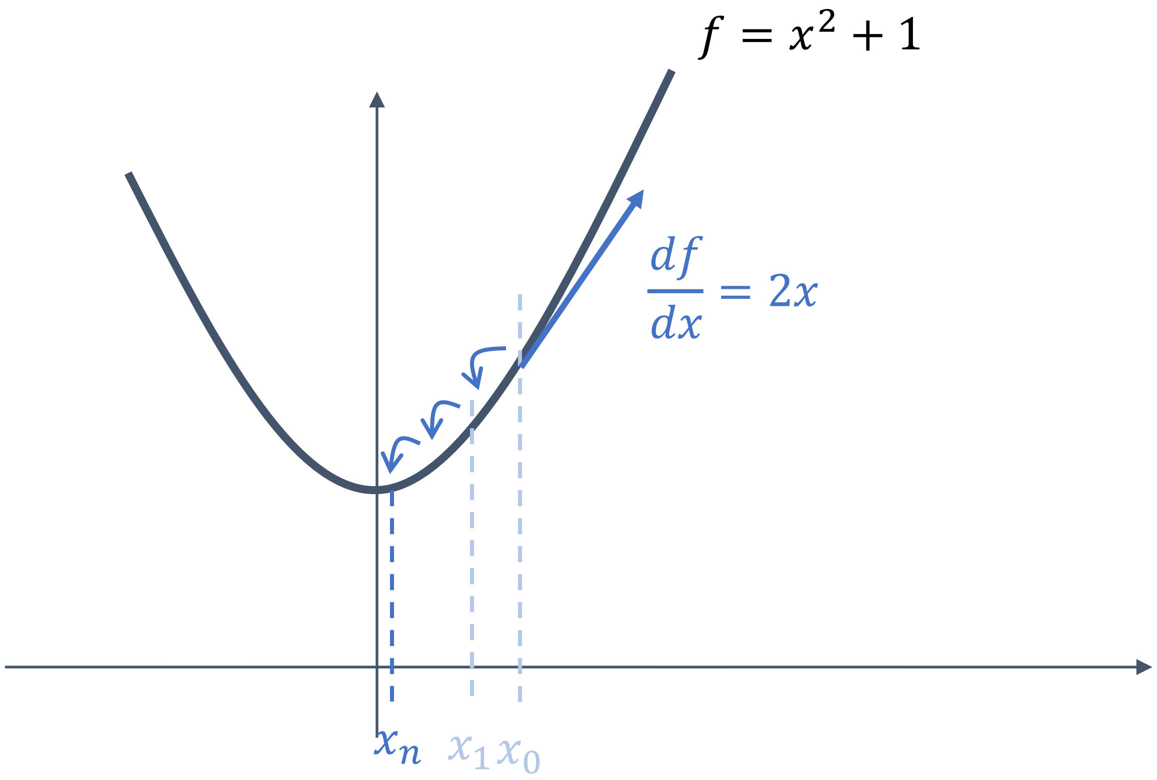

Gradient descent

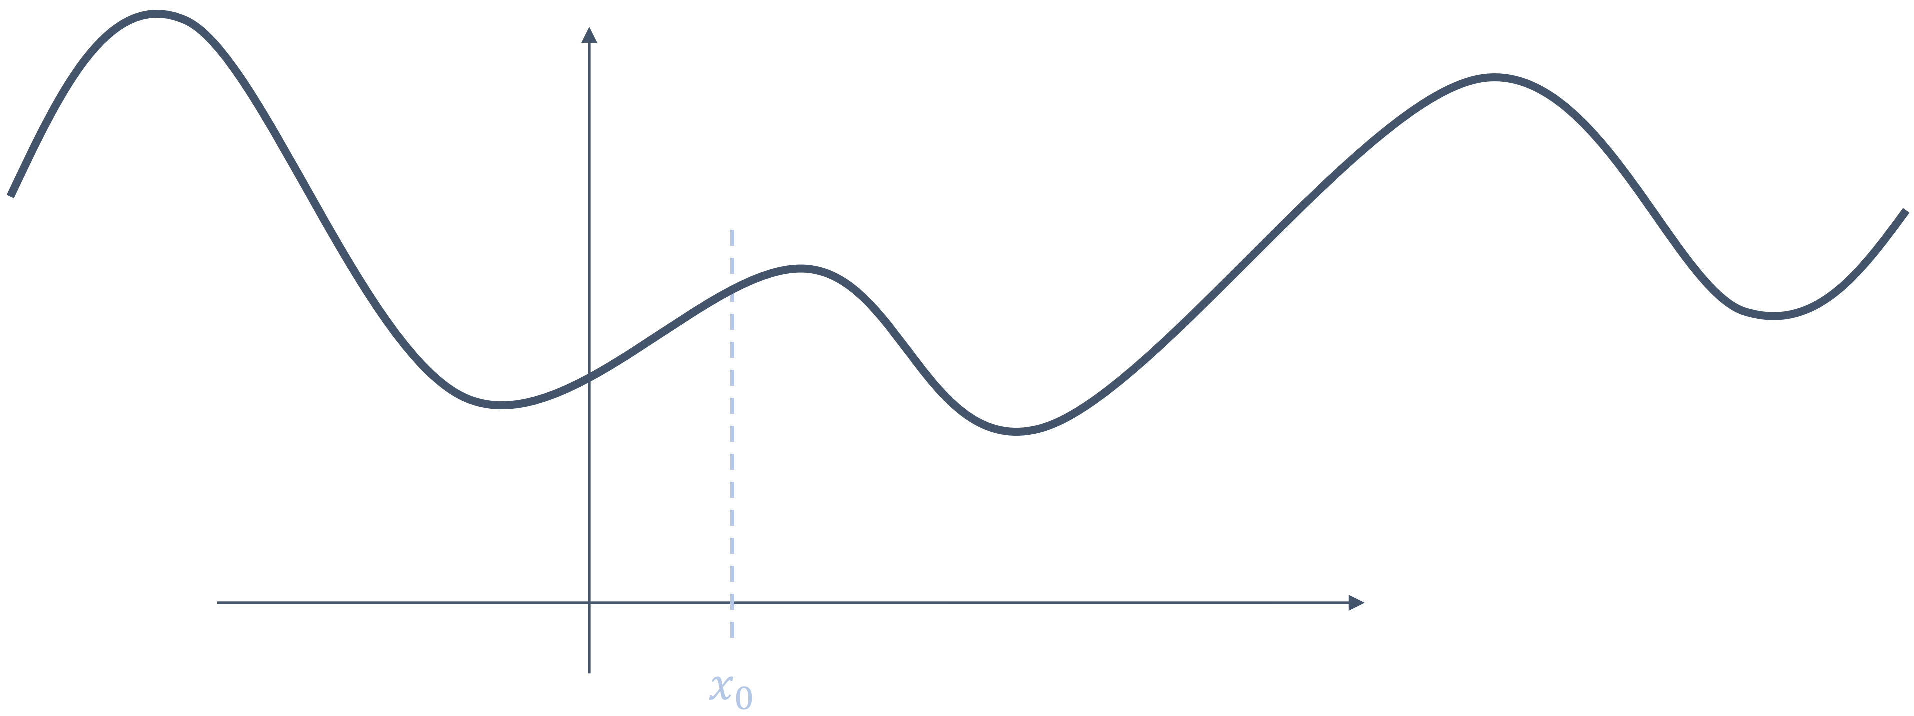

Gradient descent

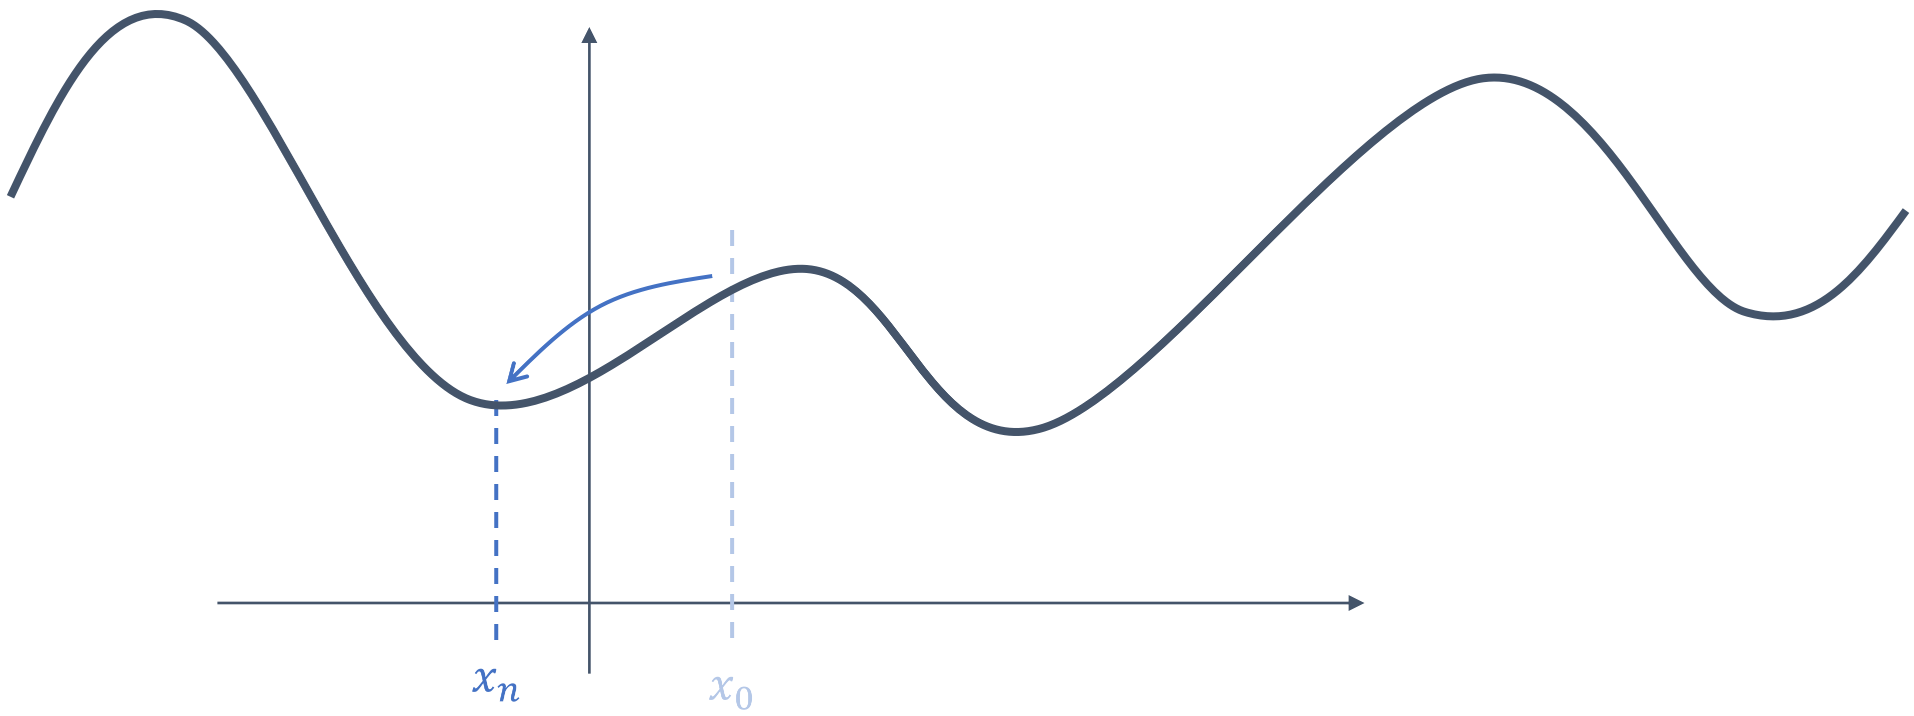

Gradient descent

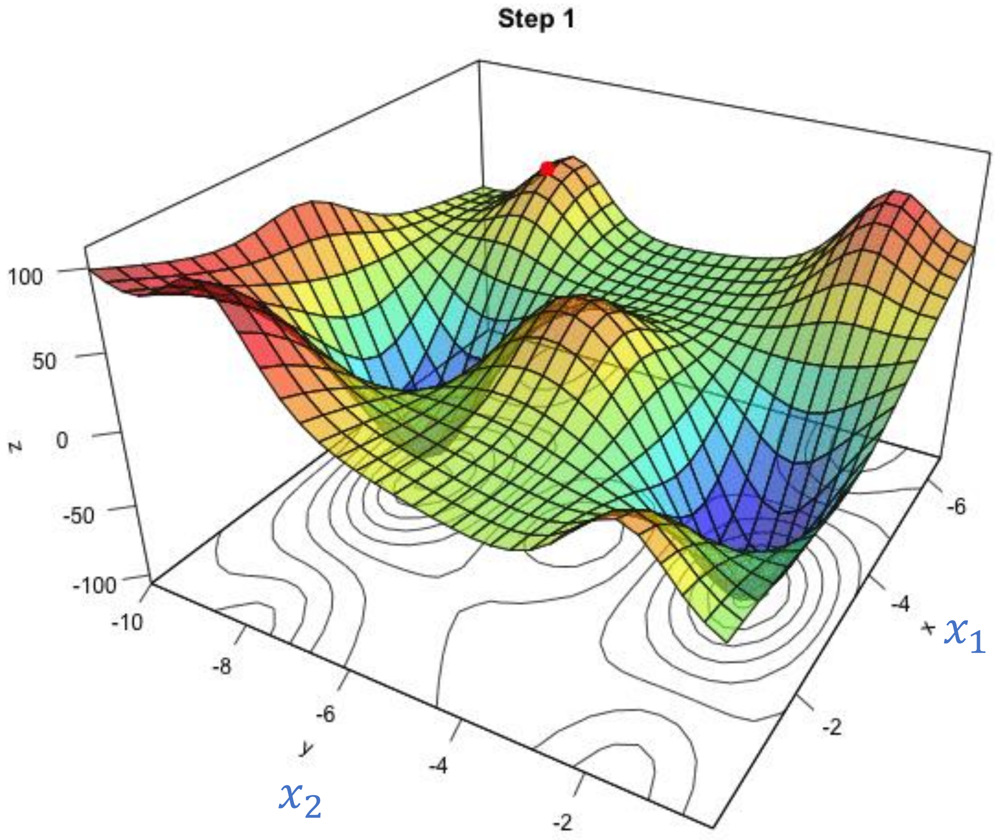

$$\begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \end{bmatrix}$$

Gradient descent

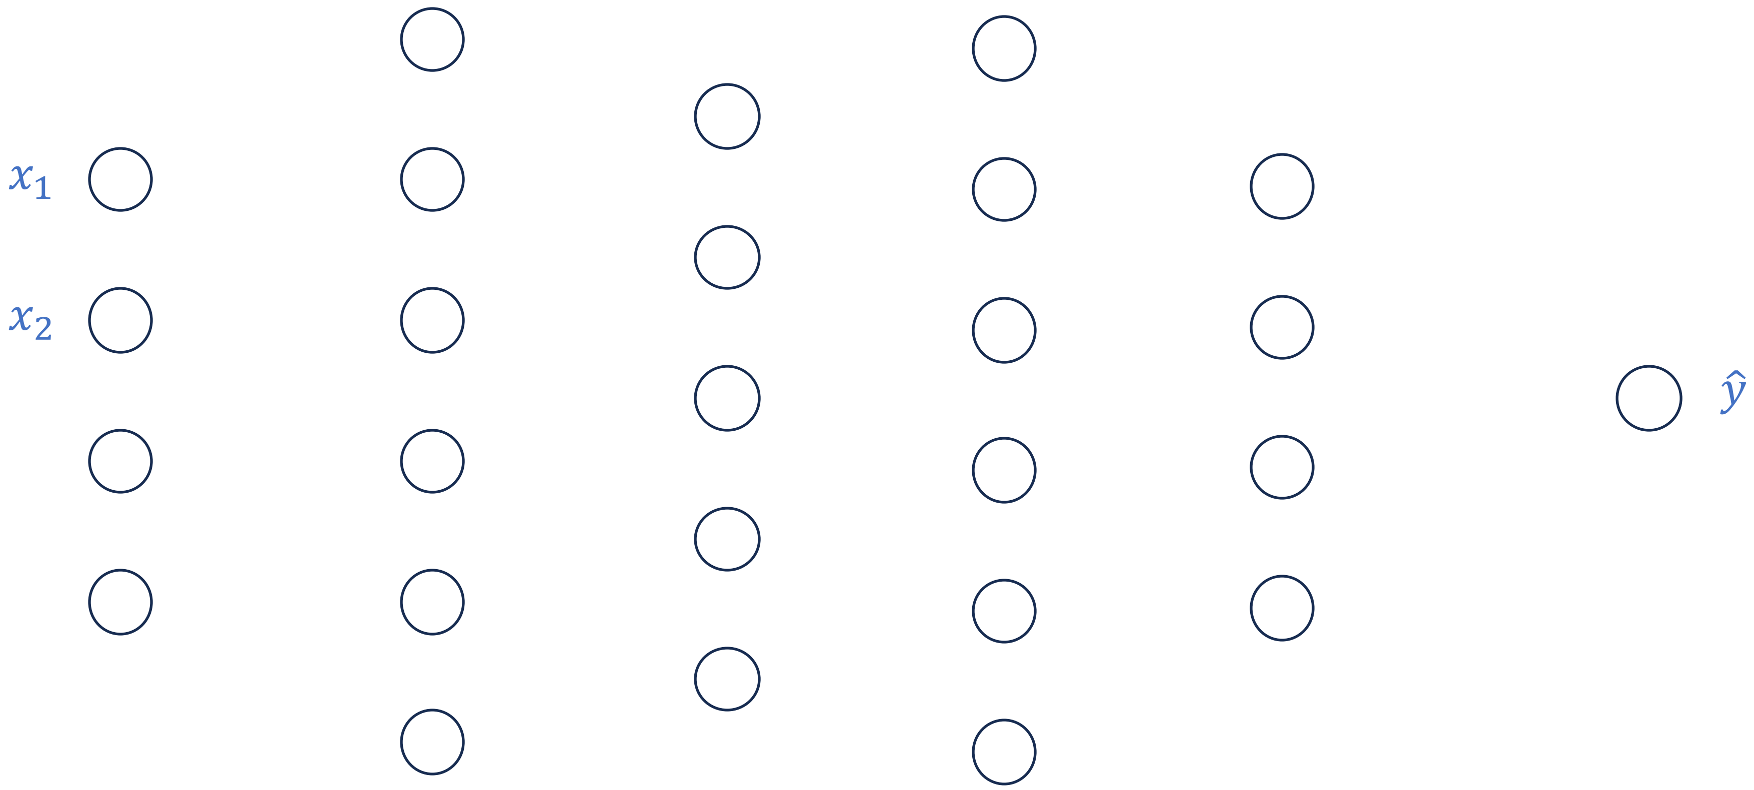

Gradient descent

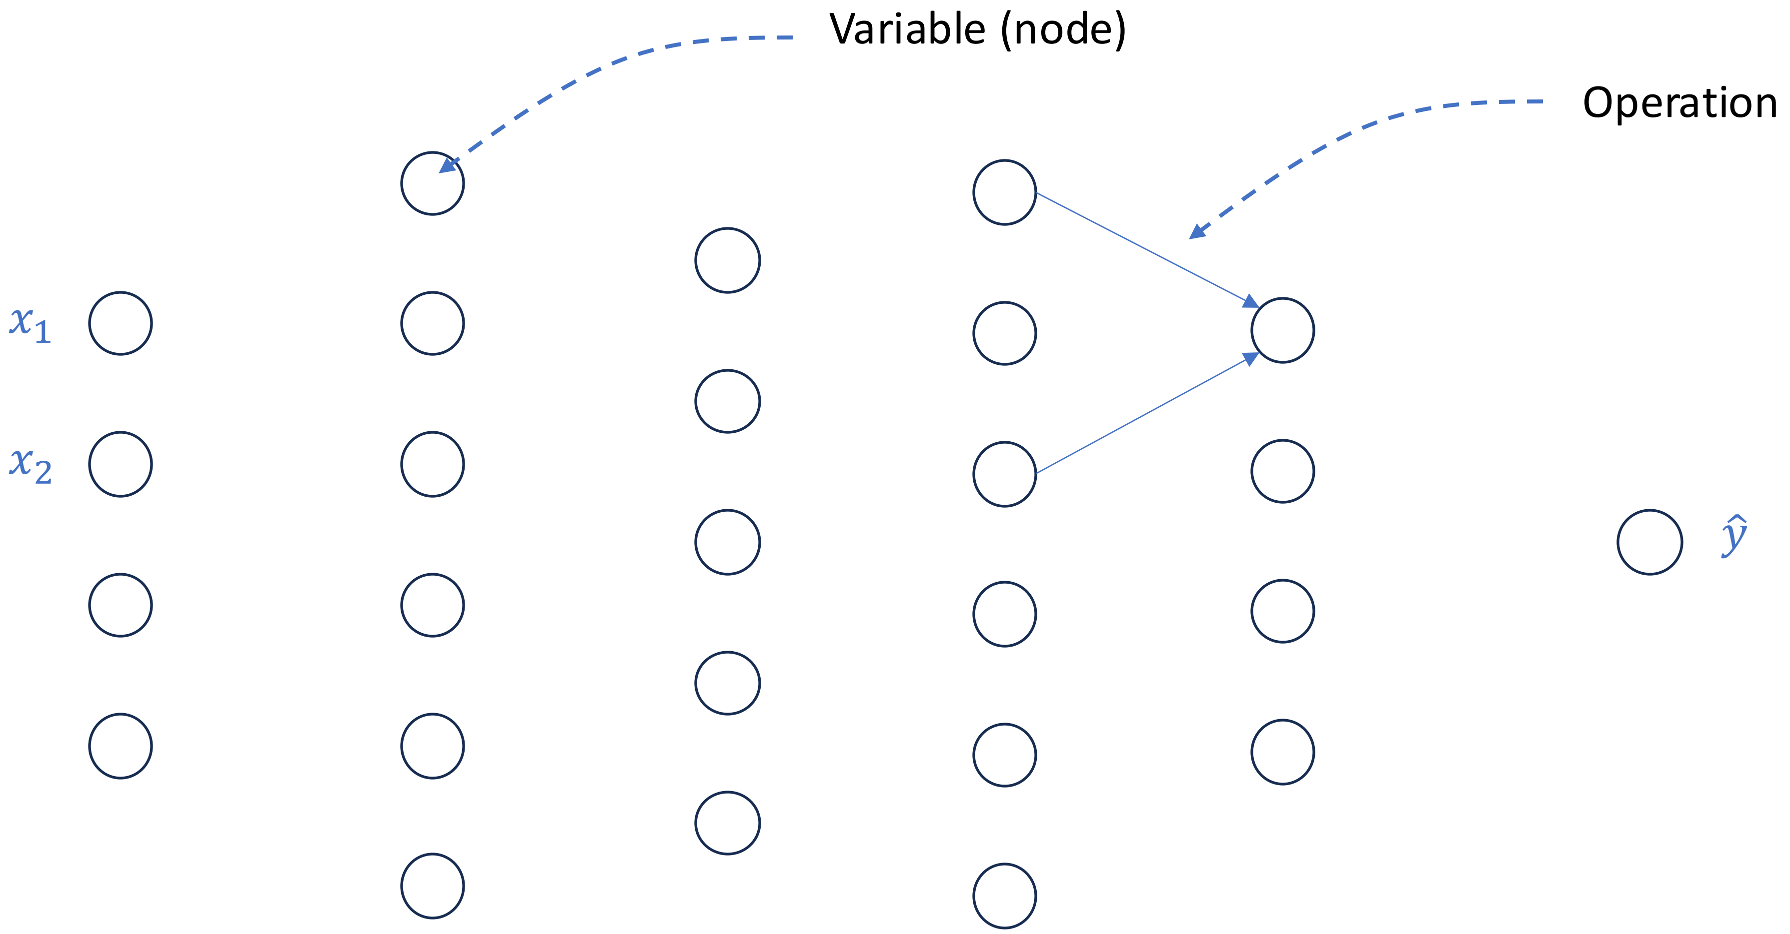

Gradient descent

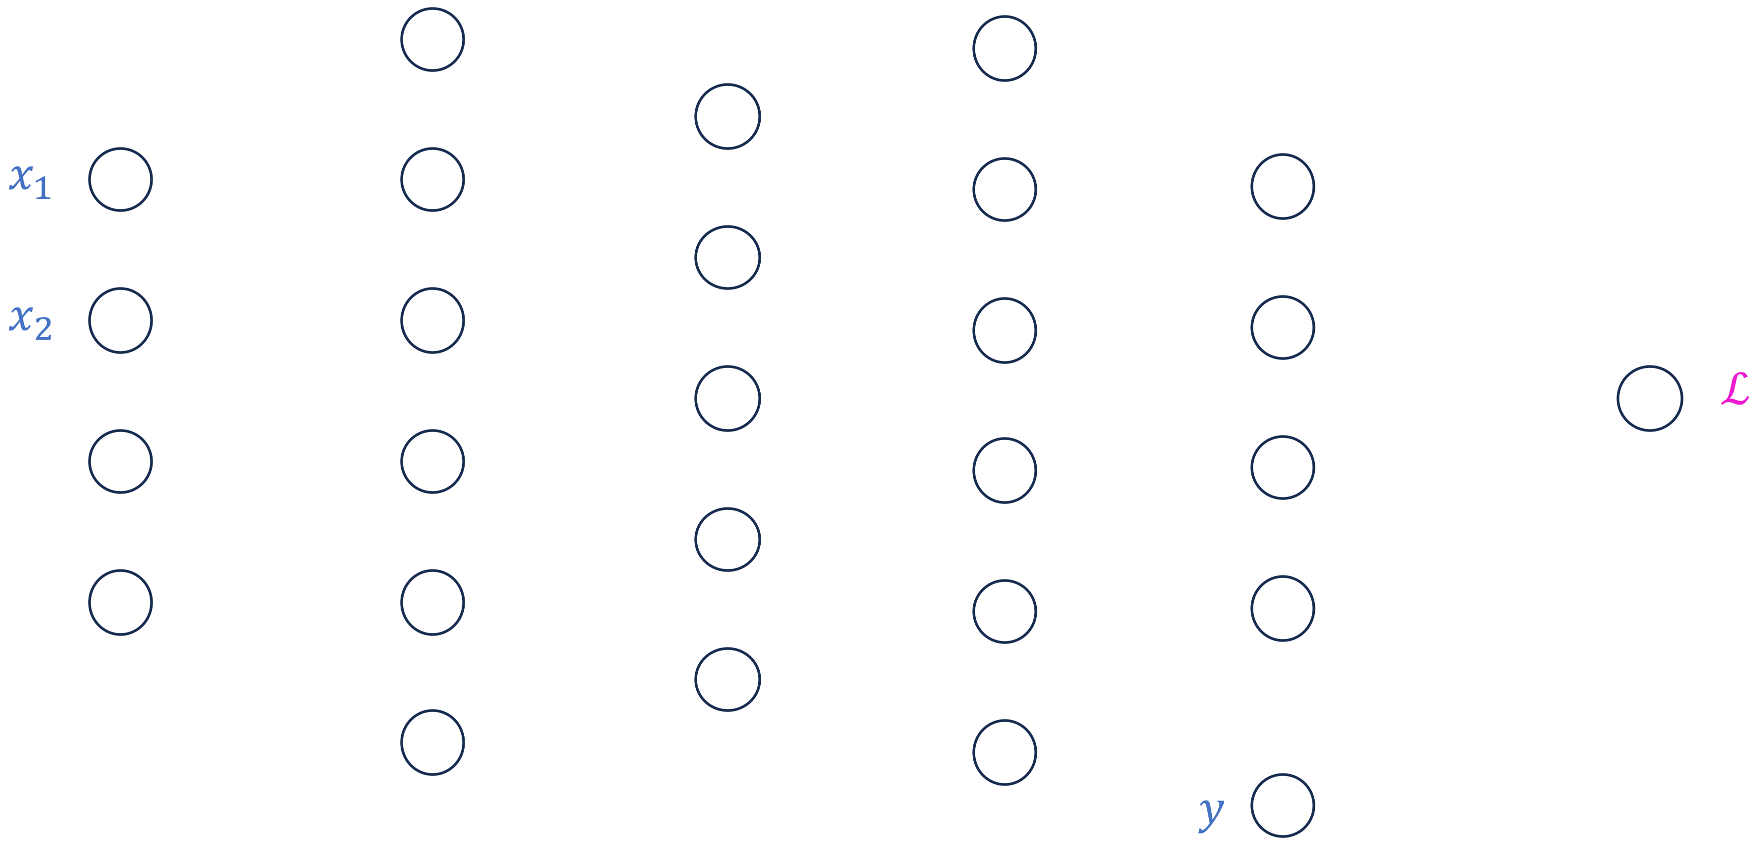

Gradient descent

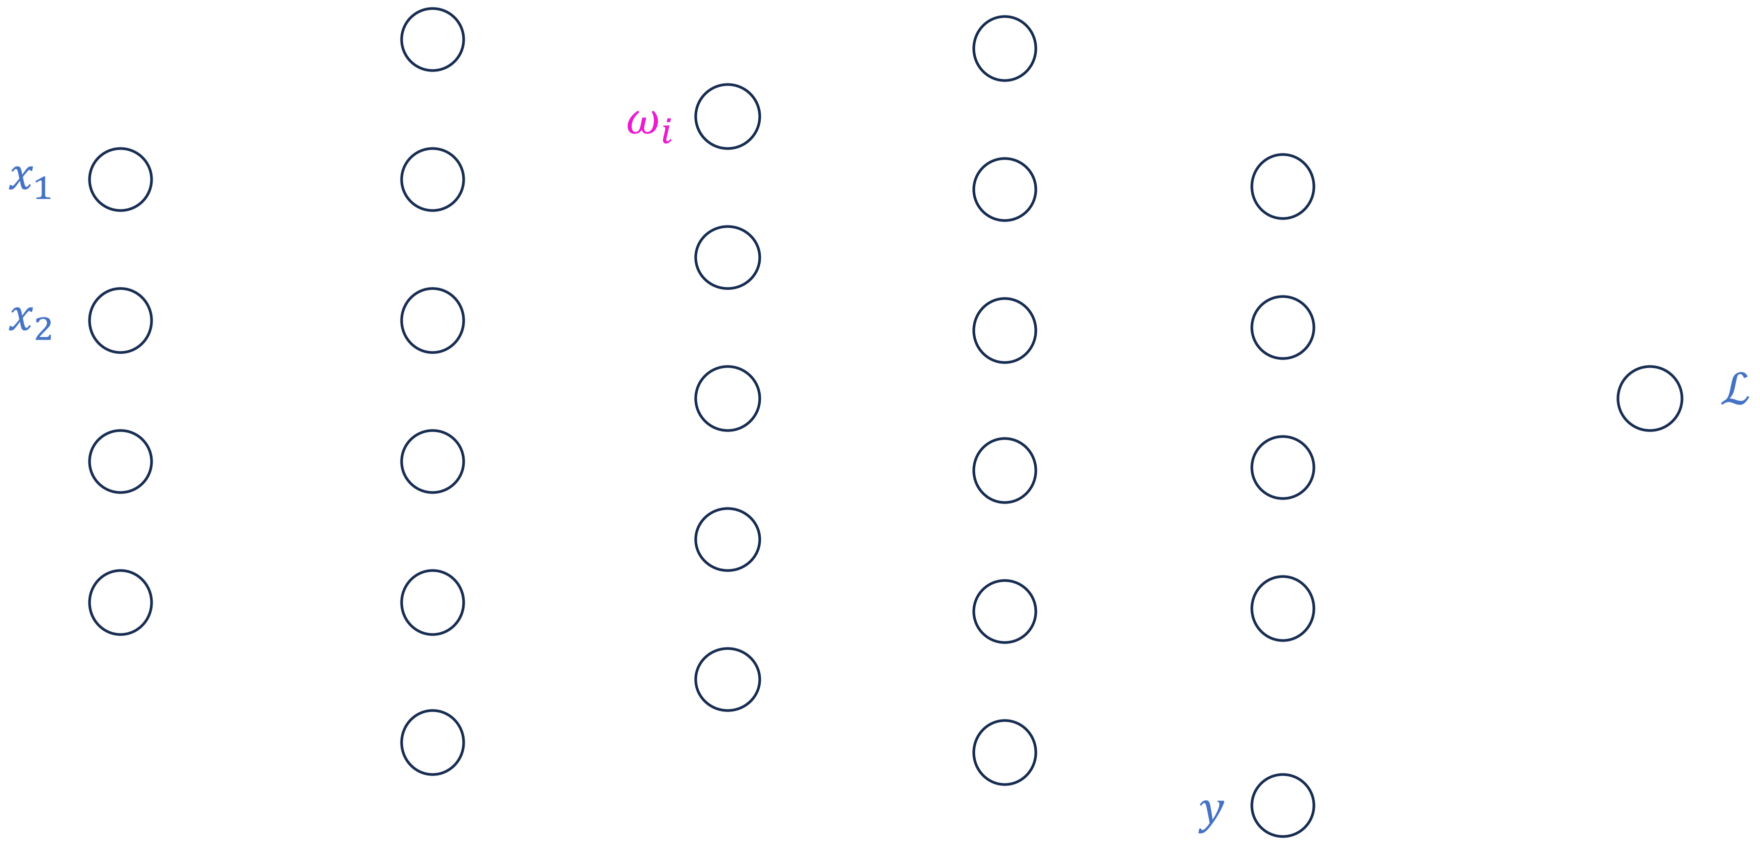

Gradient descent

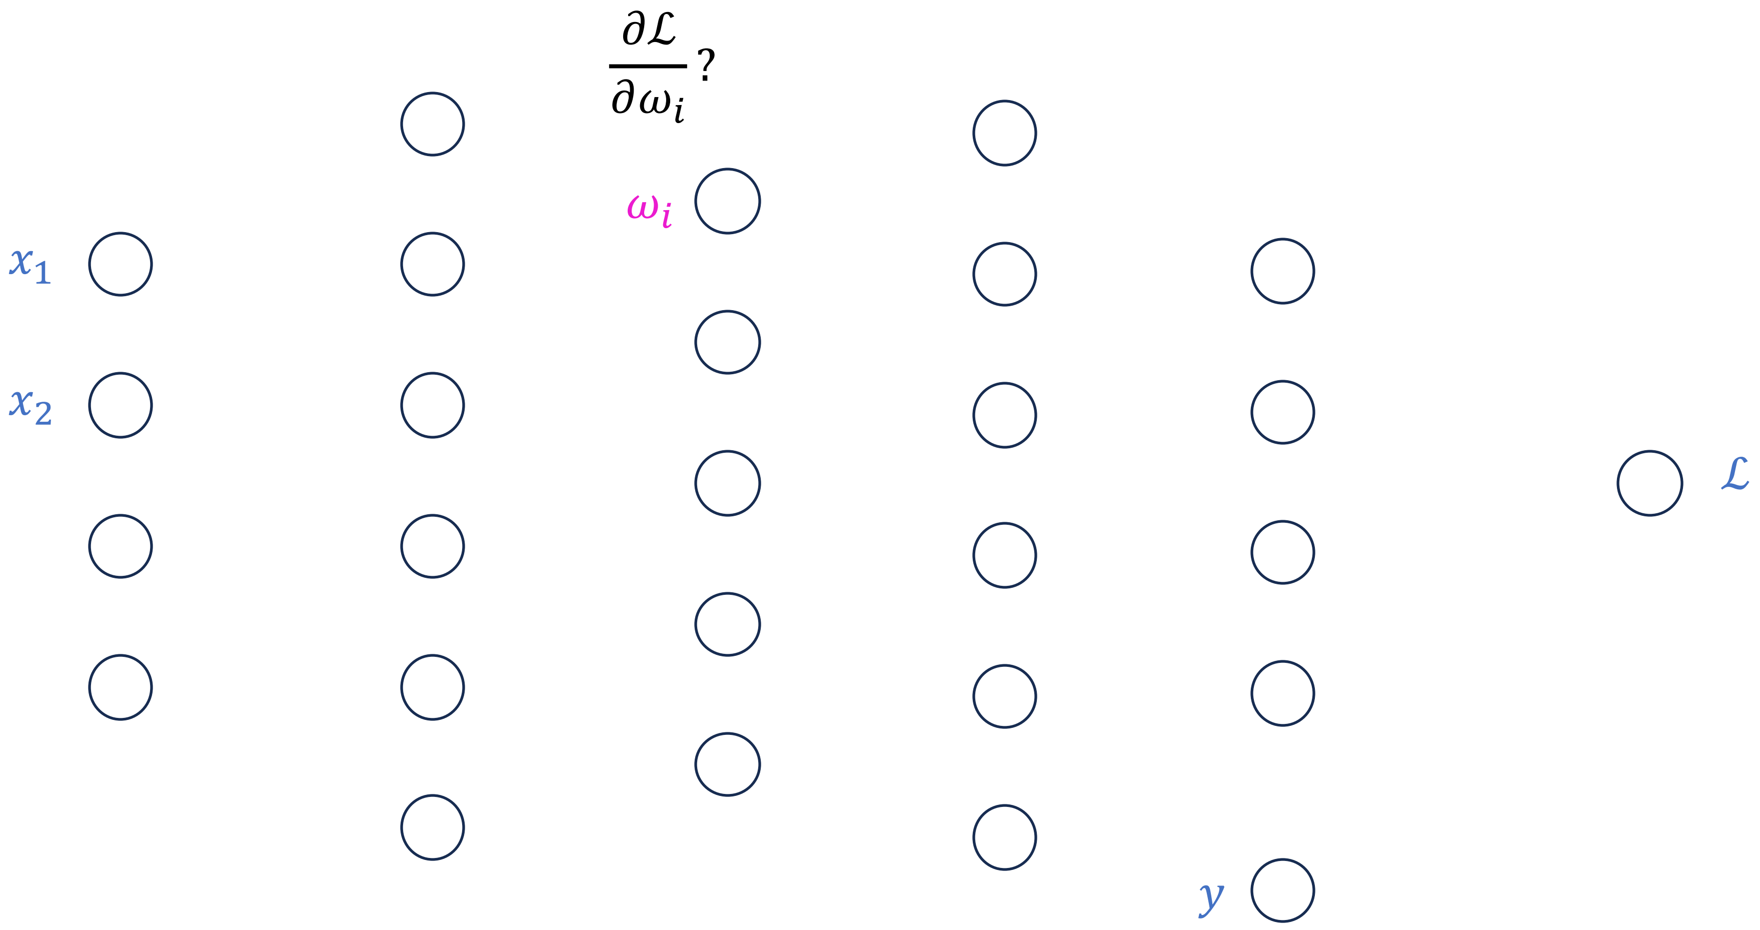

Gradient descent

$$\omega_i \mathrel{-}= \eta \frac{\partial \mathcal{L}}{\partial \omega_i}$$

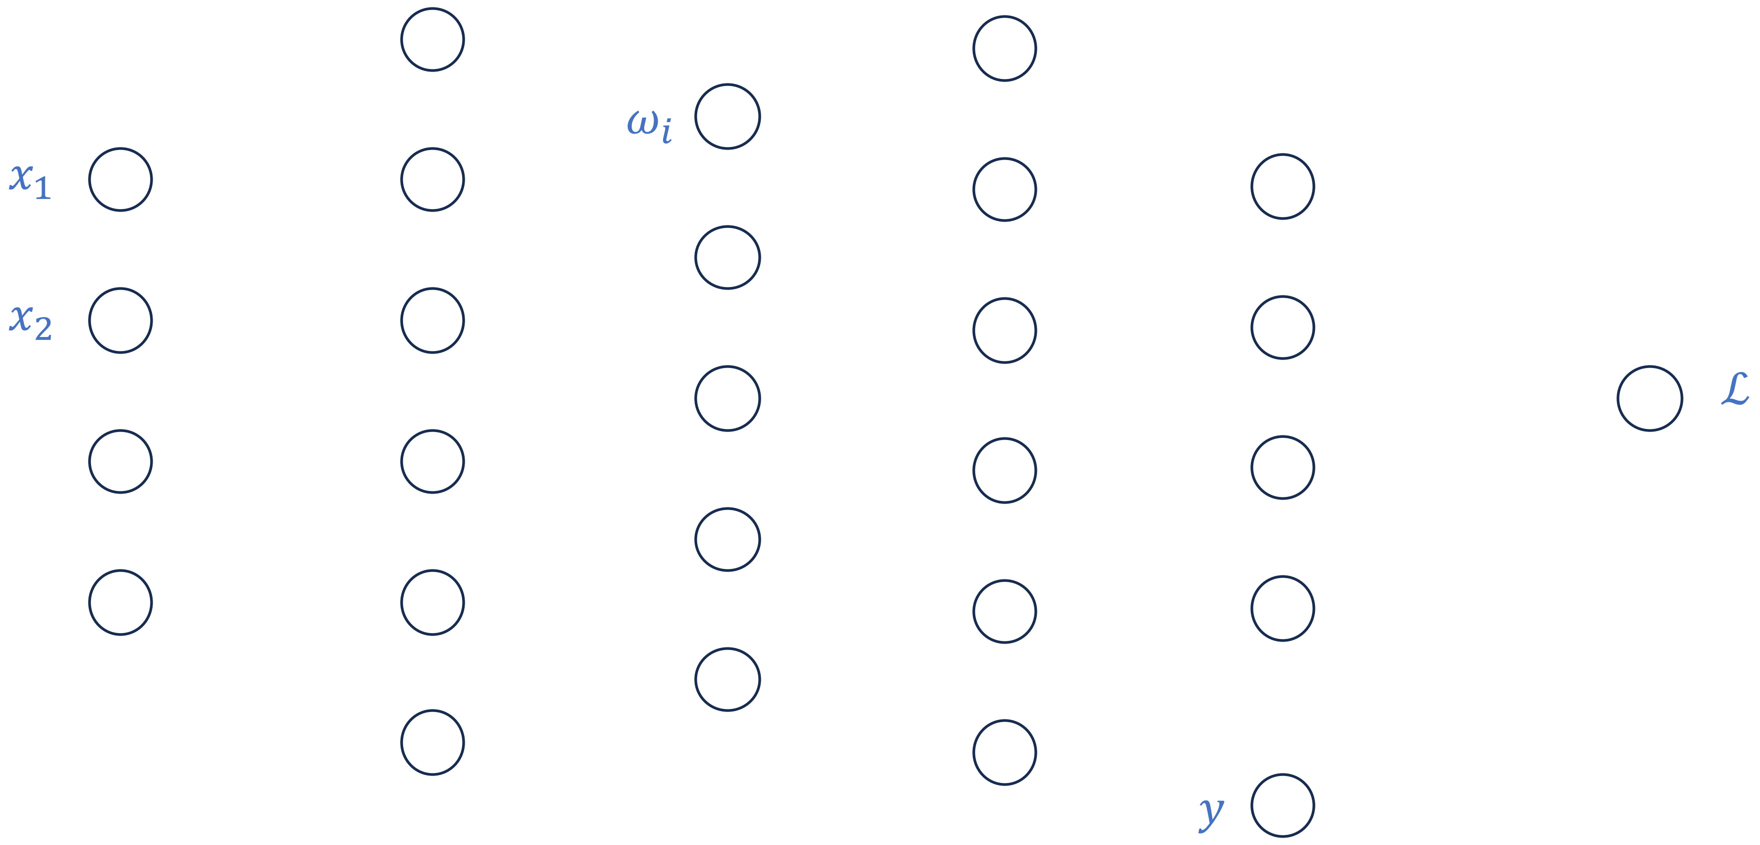

Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial a}\frac{\partial a}{\partial \omega_i}$$

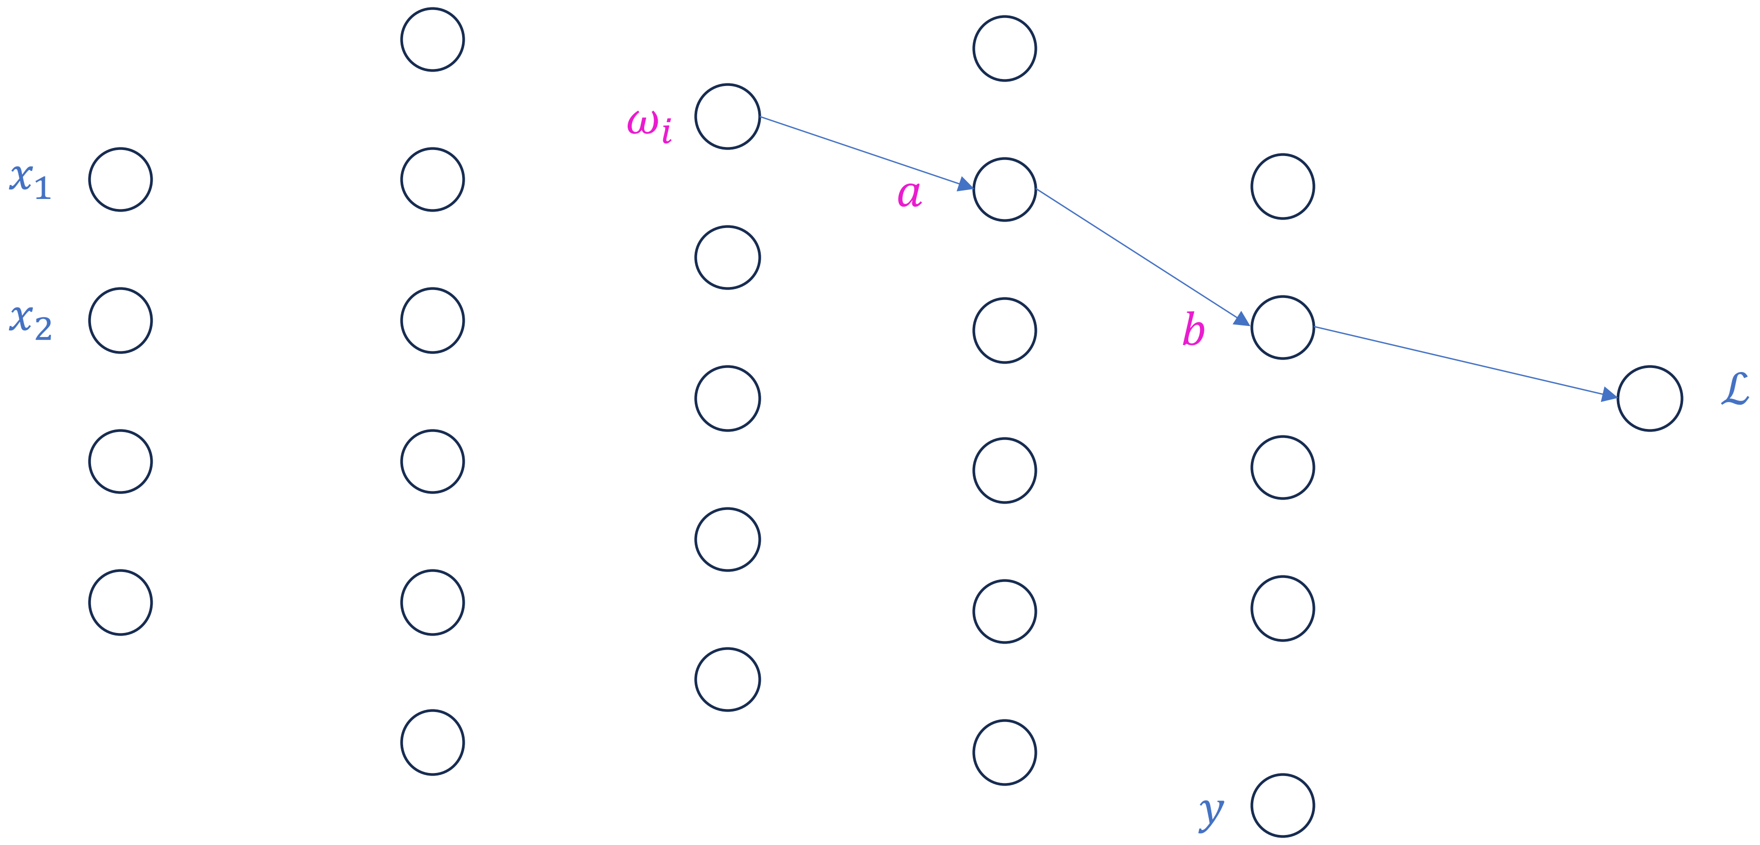

Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial a}\frac{\partial a}{\partial \omega_i}$$

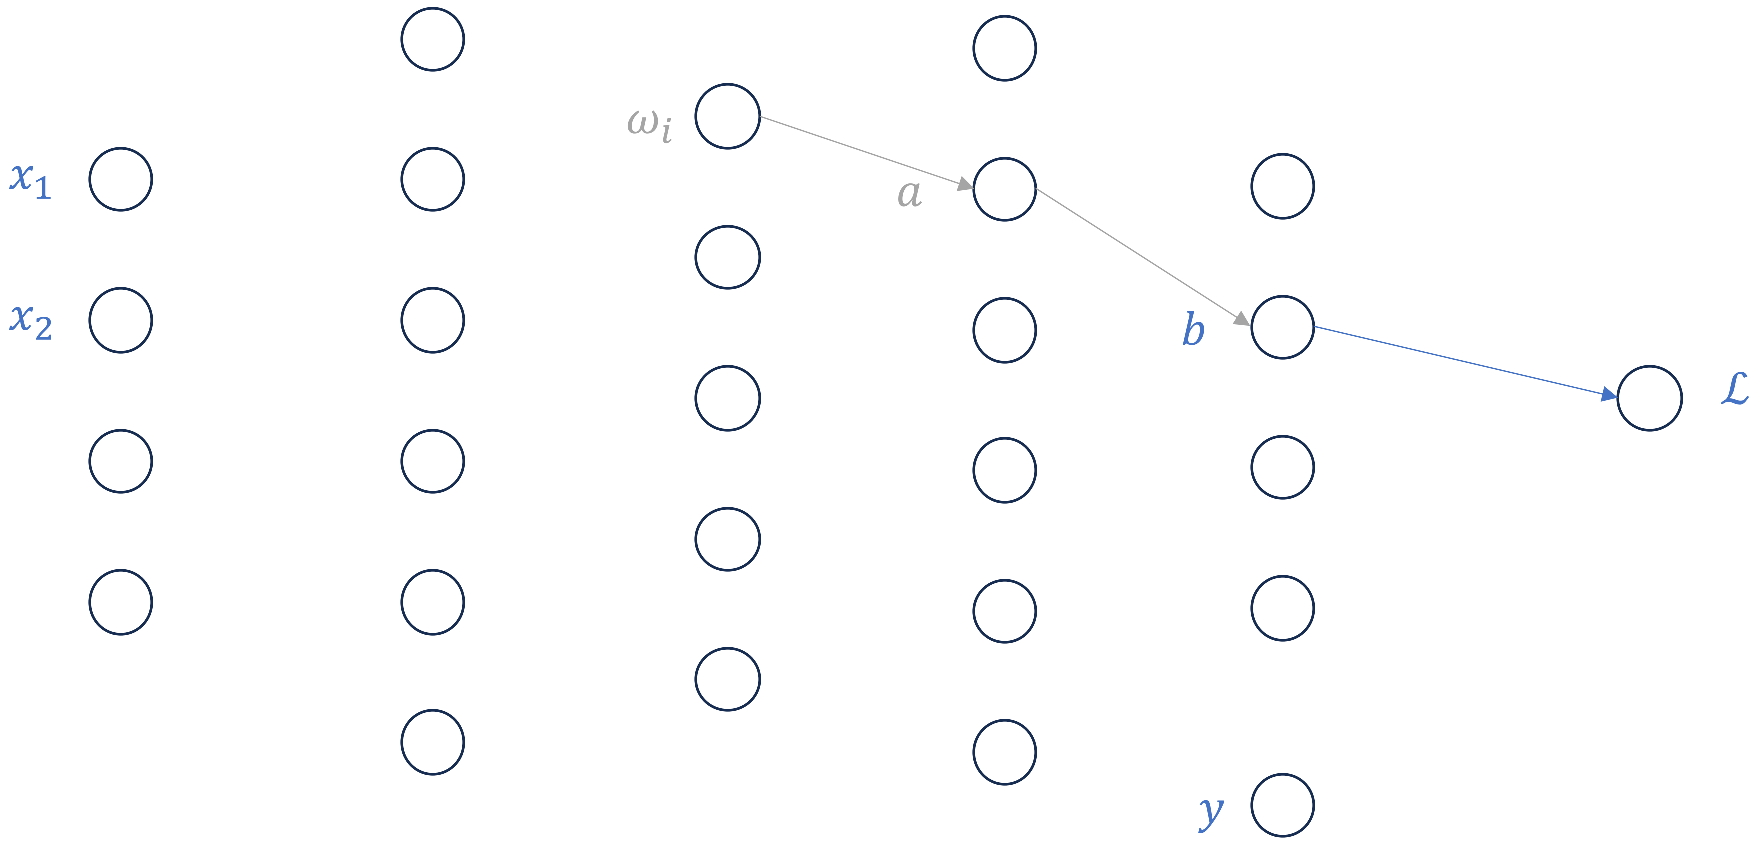

Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial a}\frac{\partial a}{\partial \omega_i}$$

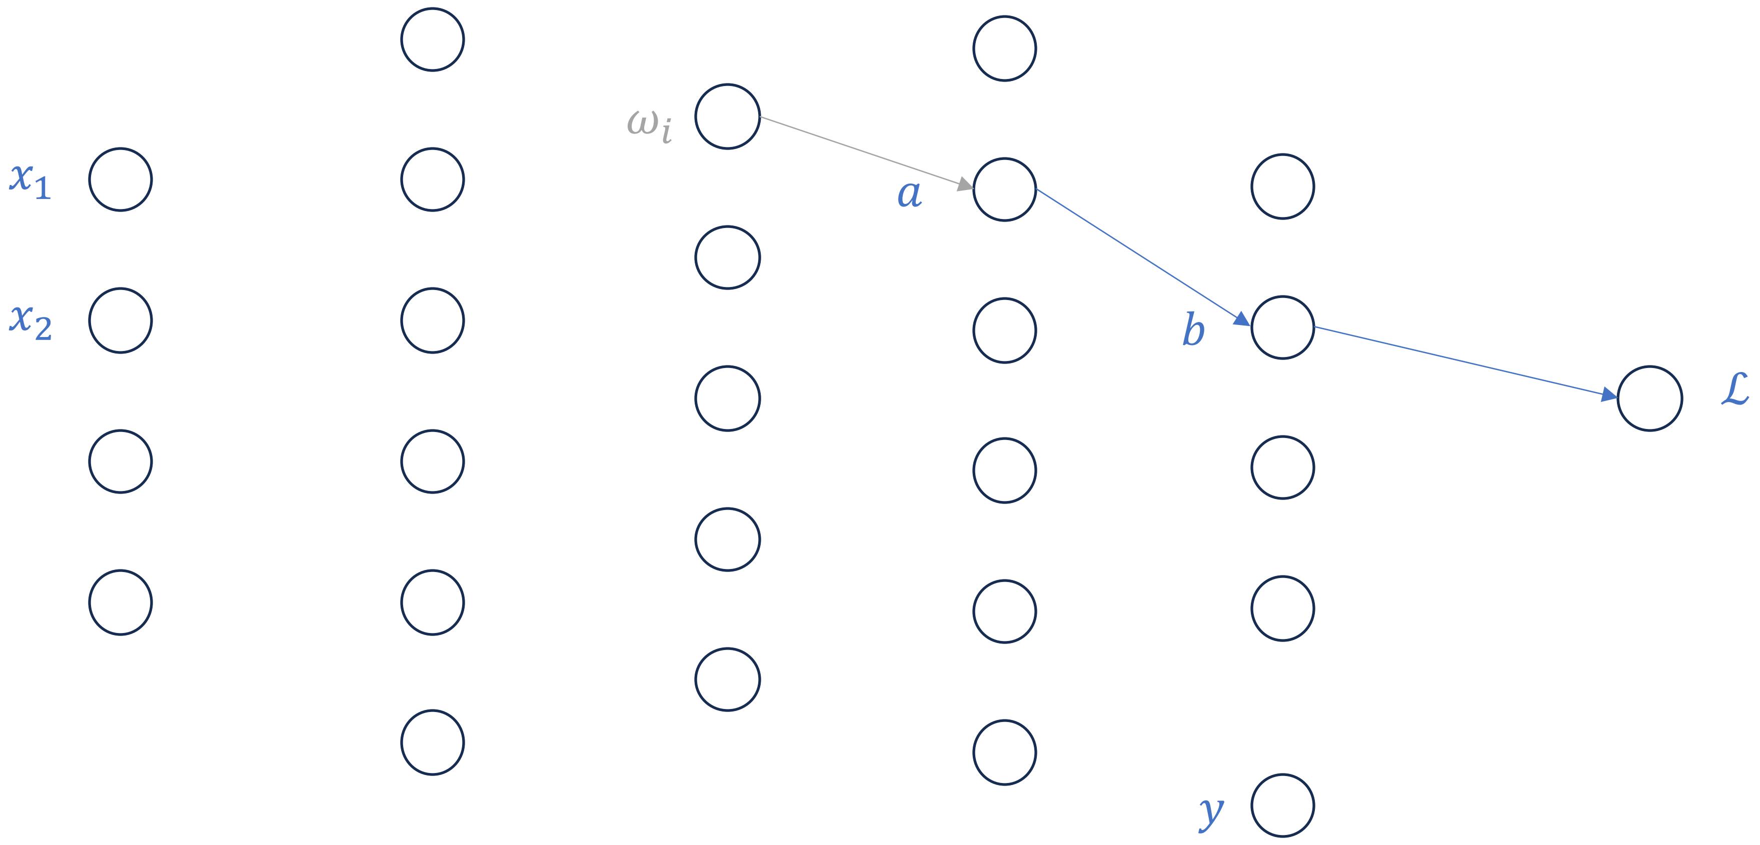

Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial a}\frac{\partial a}{\partial \omega_i}$$

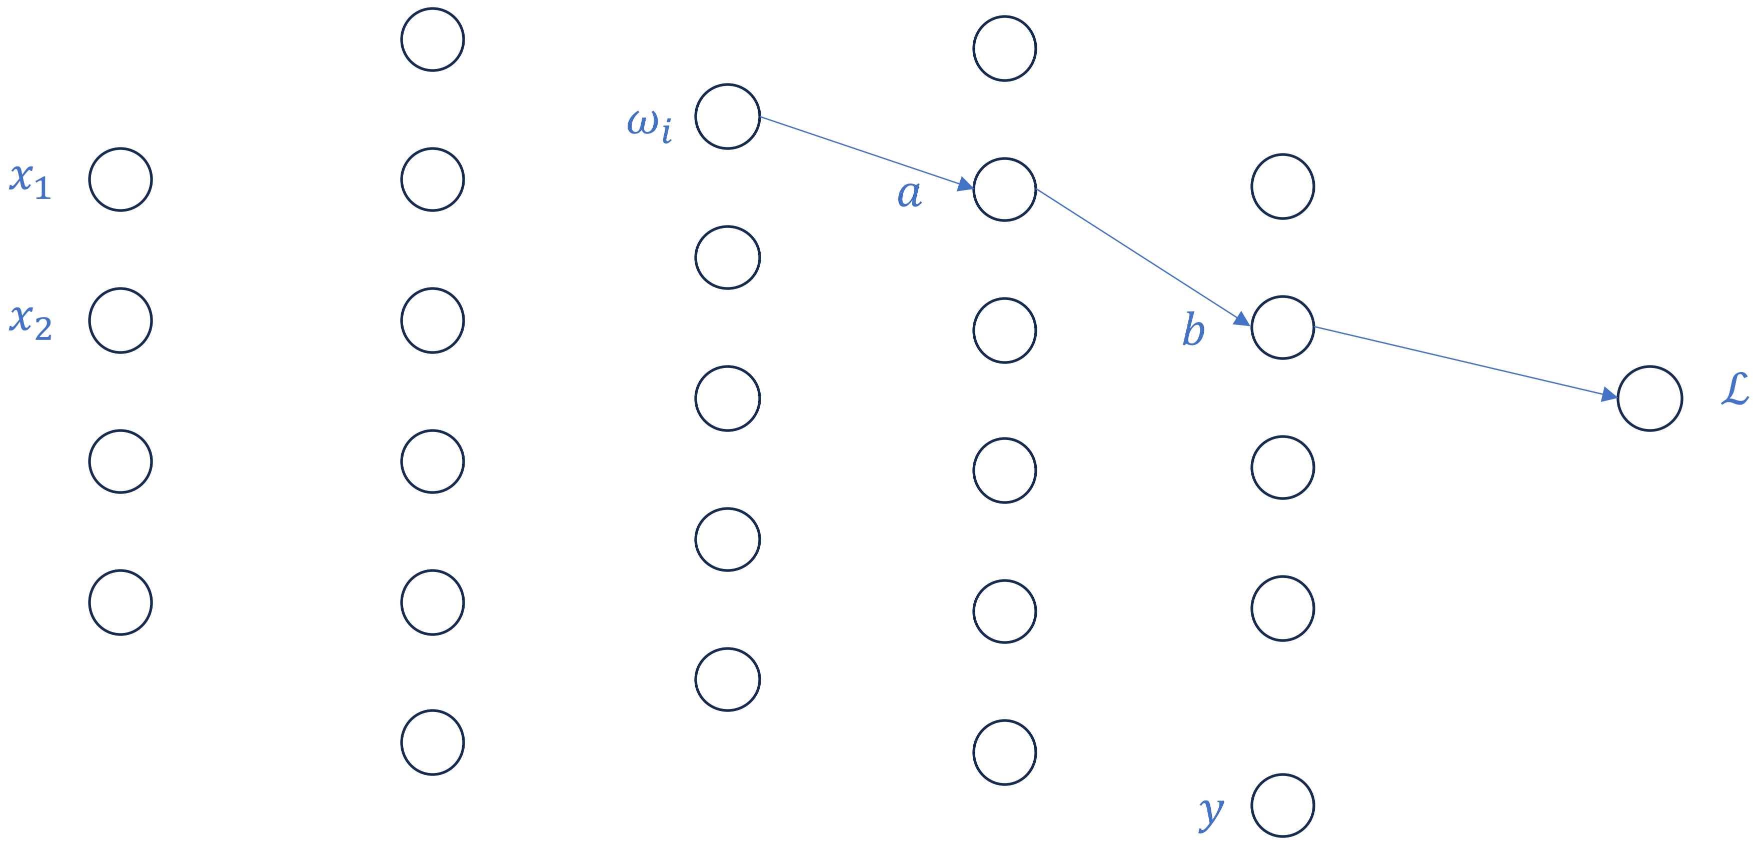

Chain rule

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial a}\frac{\partial a}{\partial \omega_i}$$

In a neural network, all these partial derivatives are easy to compute

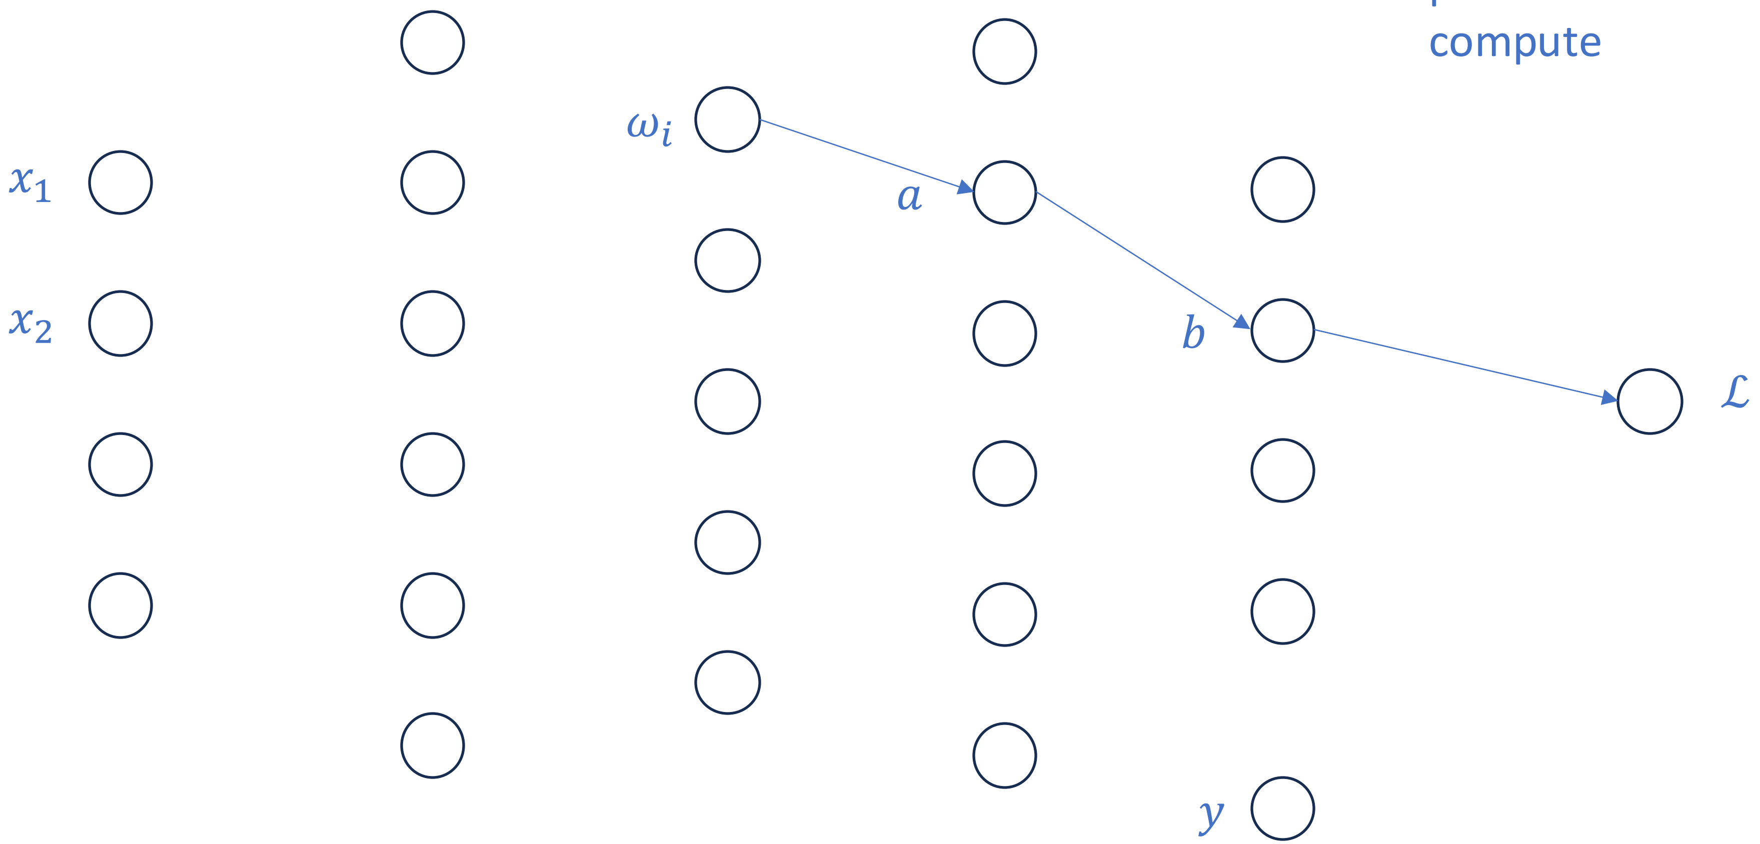

Chain rule

$$\frac{\partial \mathcal{L}}{\partial x_i} ?!?$$

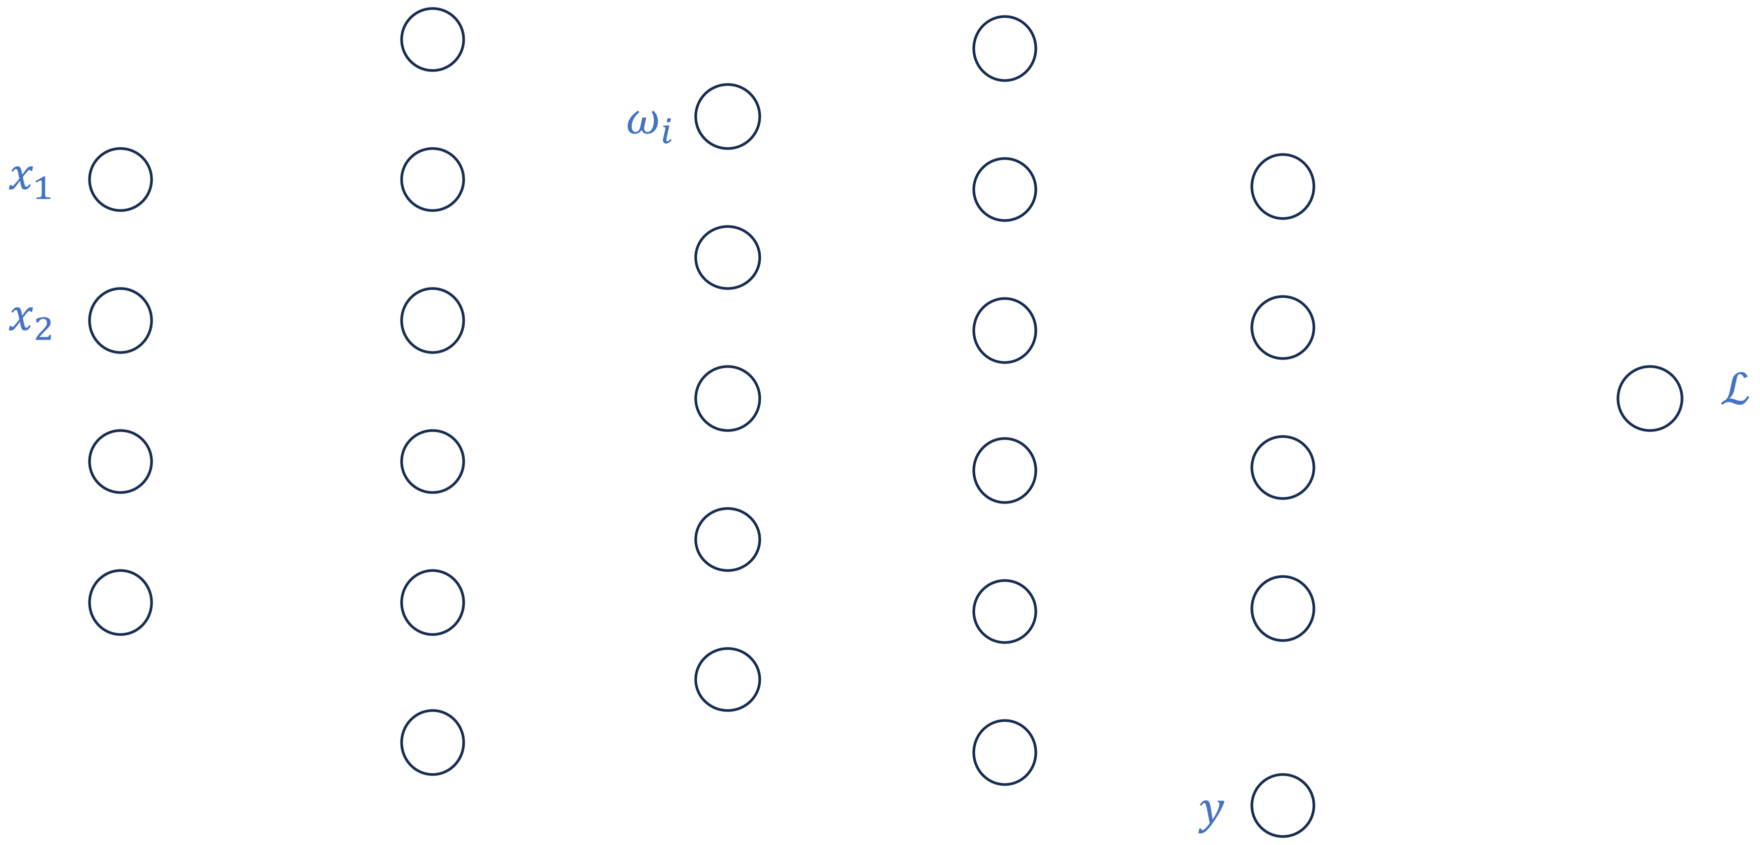

The gradient flow

$\frac{\partial \mathcal{L}}{\partial a}$ "flows" backward to e, f and g

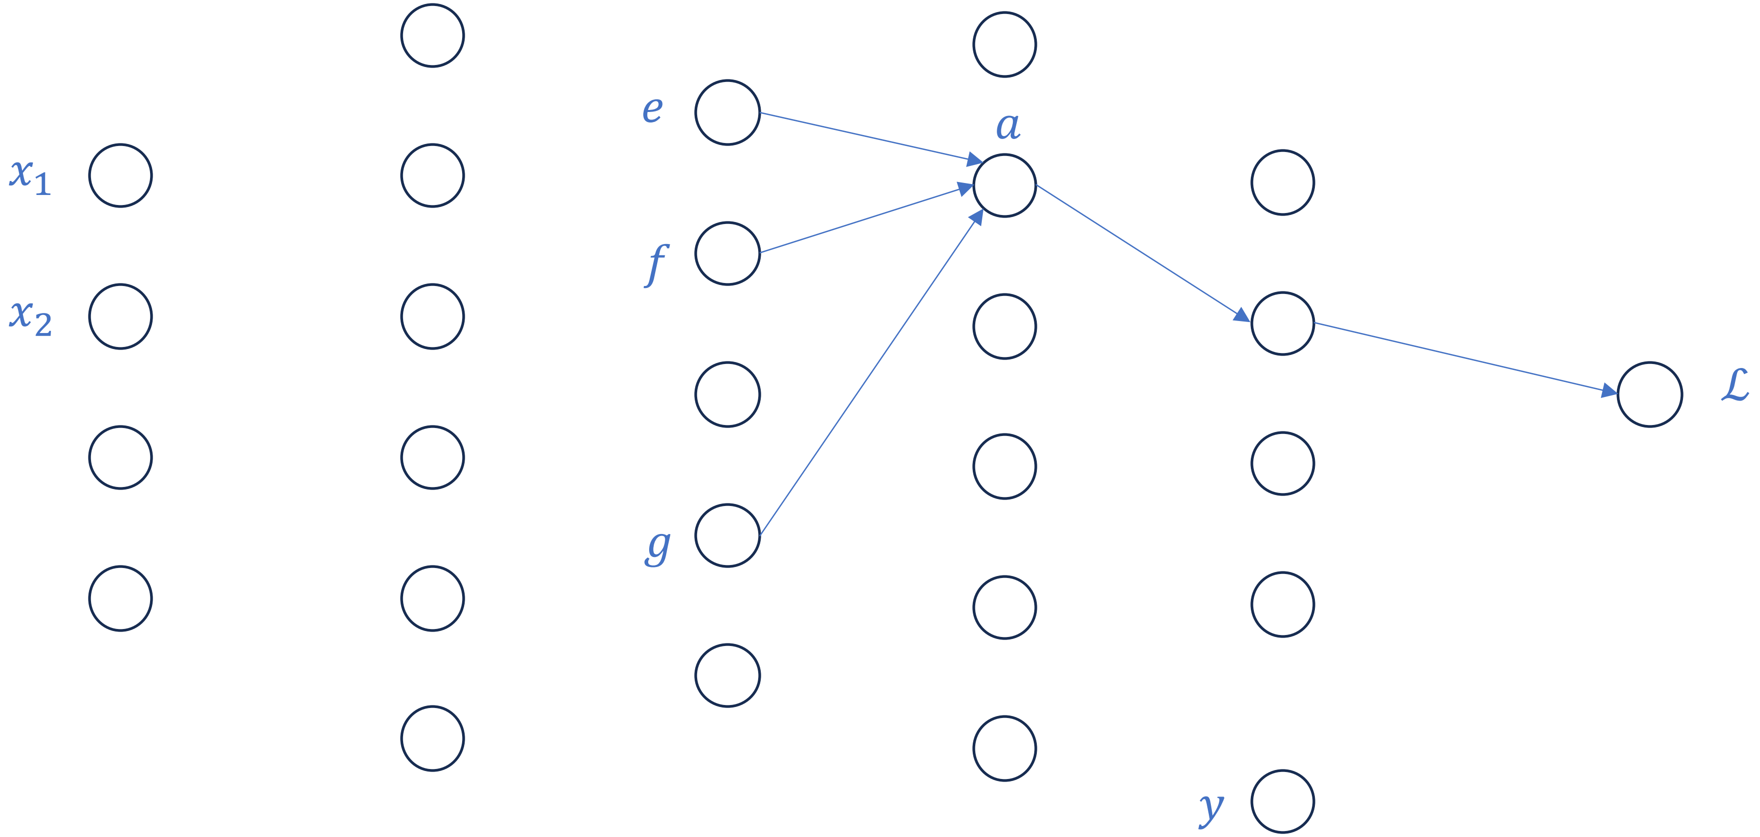

The gradient flow

Given $\frac{\partial \mathcal{L}}{\partial a}$, we can update $\frac{\partial \mathcal{L}}{\partial e}$, $\frac{\partial \mathcal{L}}{\partial f}$ and $\frac{\partial \mathcal{L}}{\partial g}$

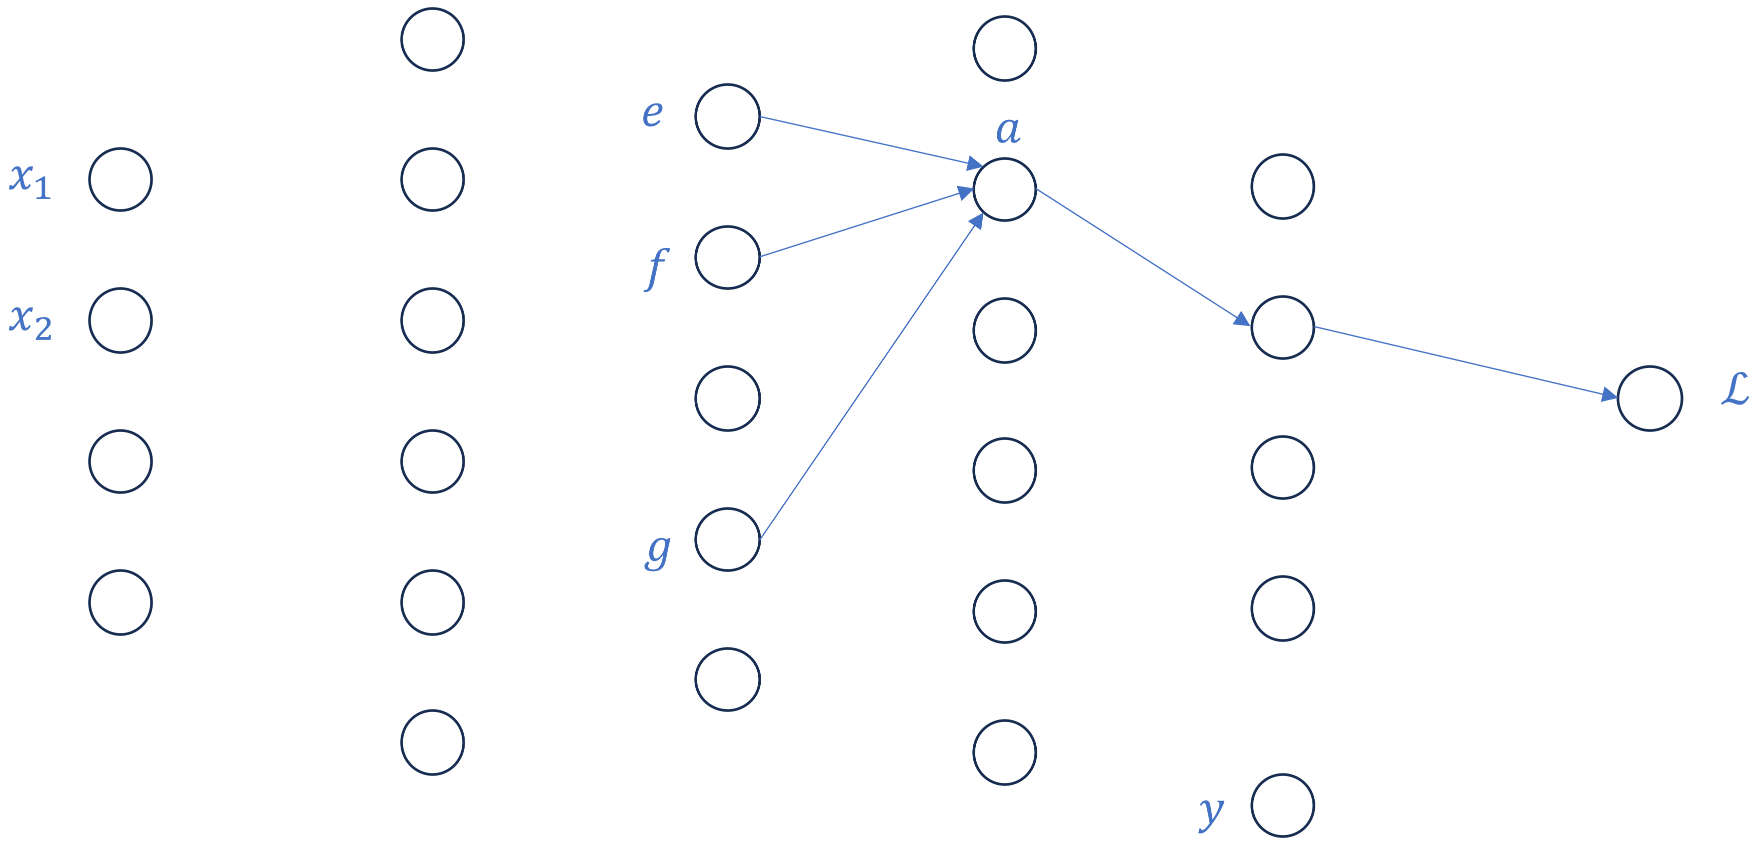

The gradient flow

$\frac{\partial \mathcal{L}}{\partial a}$ and $\frac{\partial \mathcal{L}}{\partial b}$ contribute to $\frac{\partial \mathcal{L}}{\partial f}$

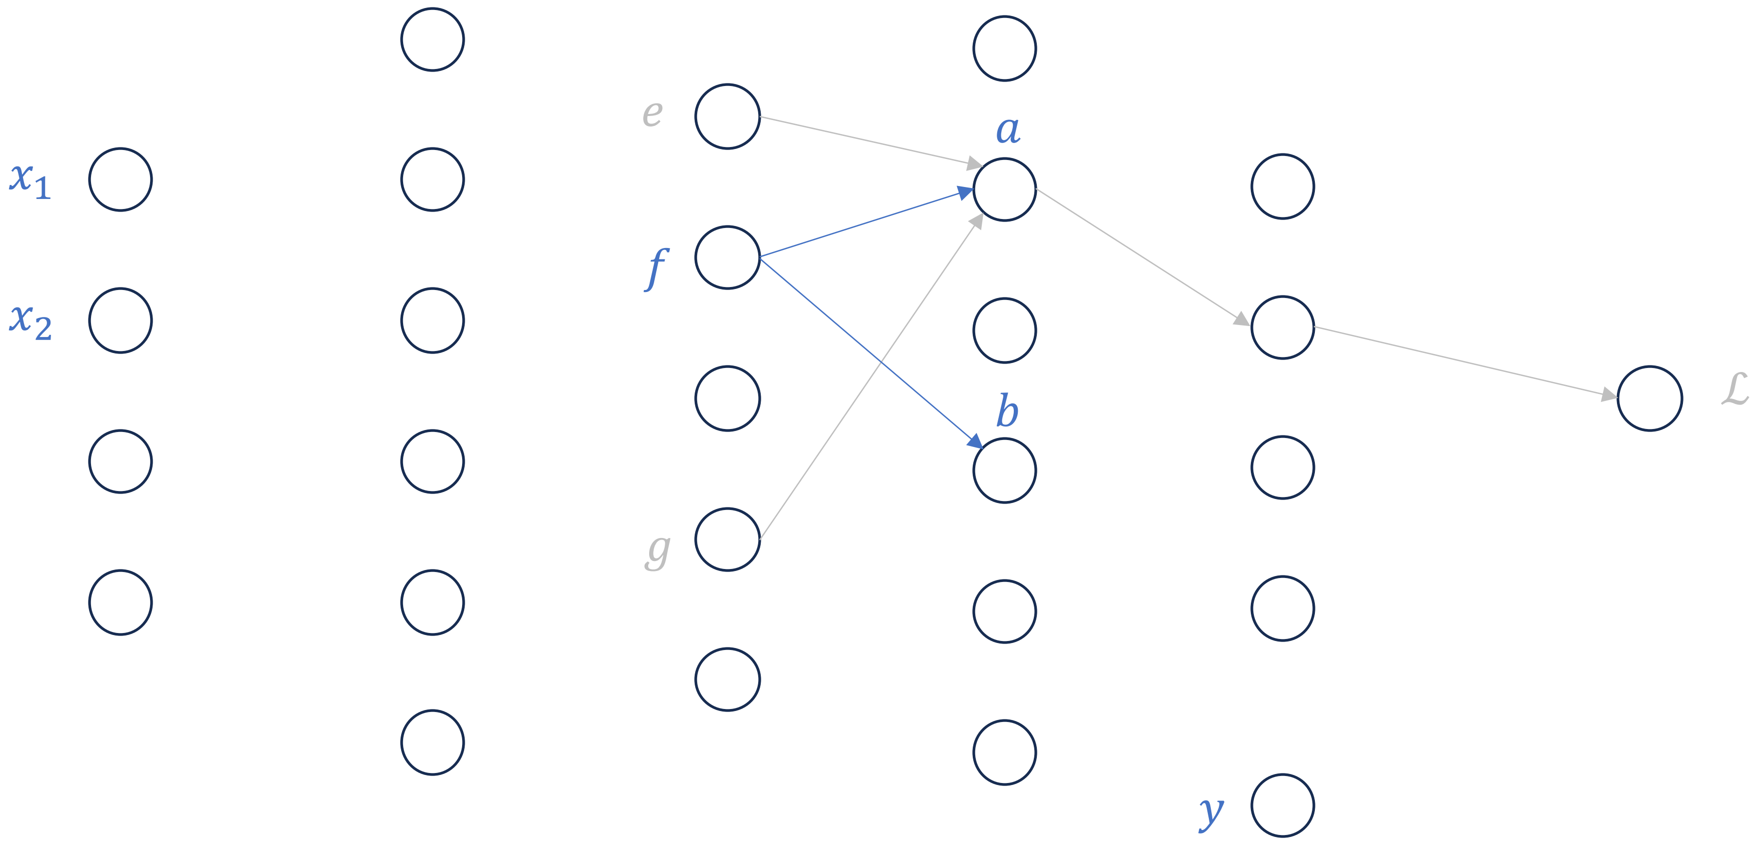

The gradient flow

$$\frac{\partial \mathcal{L}}{\partial f} = \frac{\partial \mathcal{L}}{\partial a}\frac{\partial a}{\partial f} + \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial f}$$

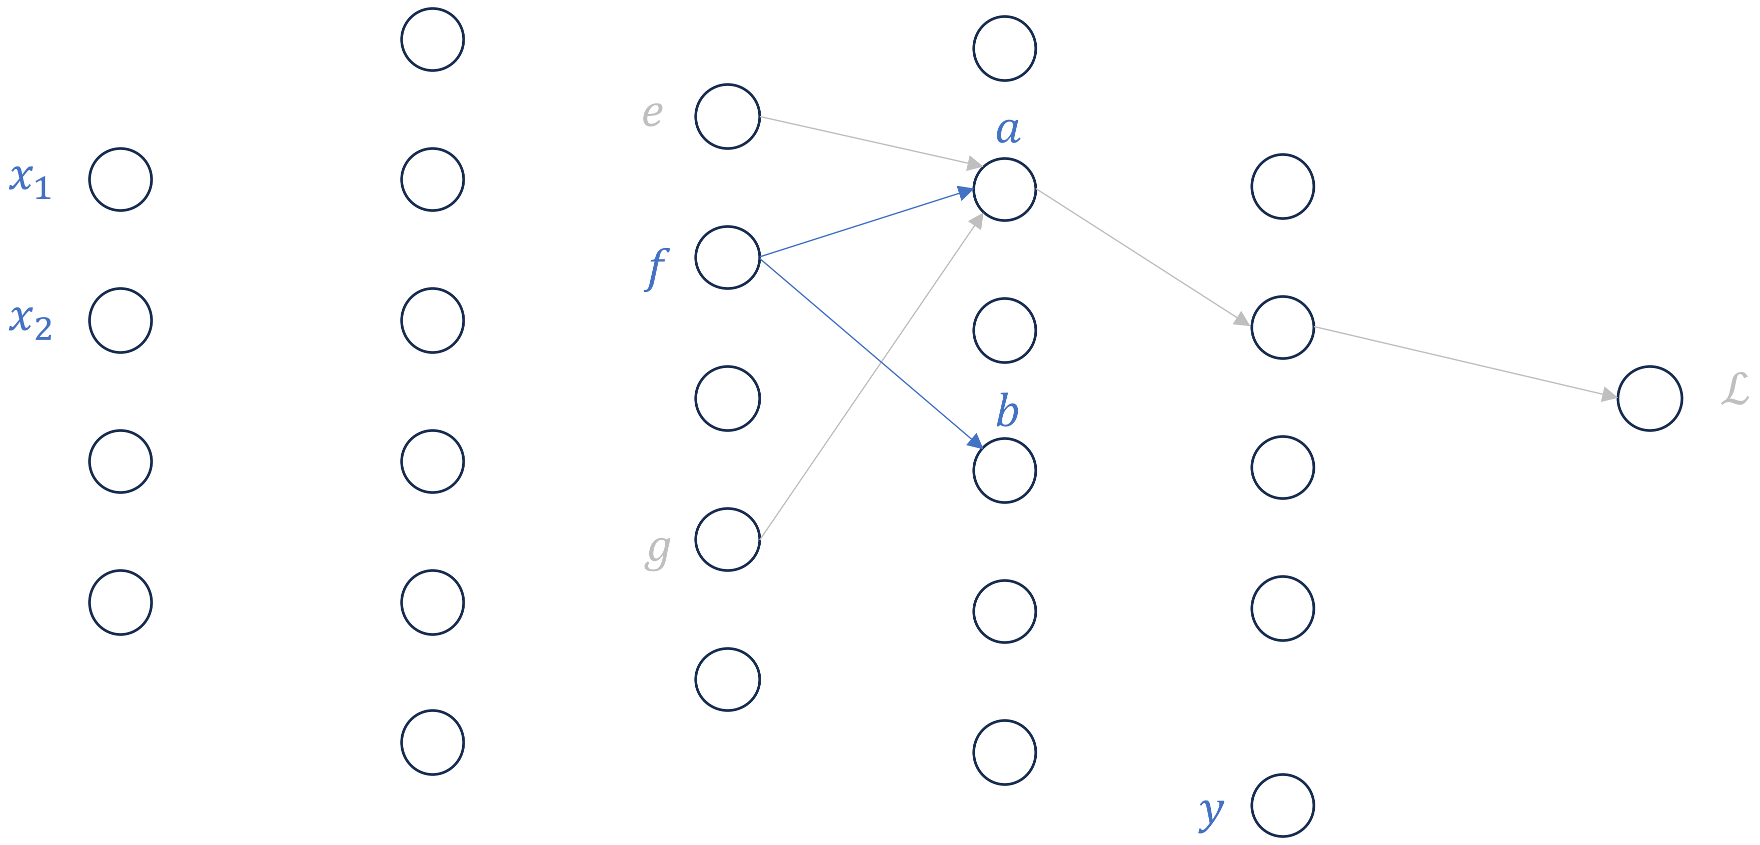

$$\mathcal{L} = f_\omega(x)$$

Assuming f_ω is smooth and differentiable w.r.t ω

$\omega = \omega_0$
Repeat:
  Compute $\mathcal{L}$
  Reset gradients $\Delta\mathcal{L} = 0$
  Compute $\Delta\mathcal{L}$ (i.e. $\partial\mathcal{L}/\partial\omega_i$)
  $\omega_i = \omega_i - \eta . \partial\mathcal{L}/\partial\omega_i$ η: learning rate

Let's build it by hand!

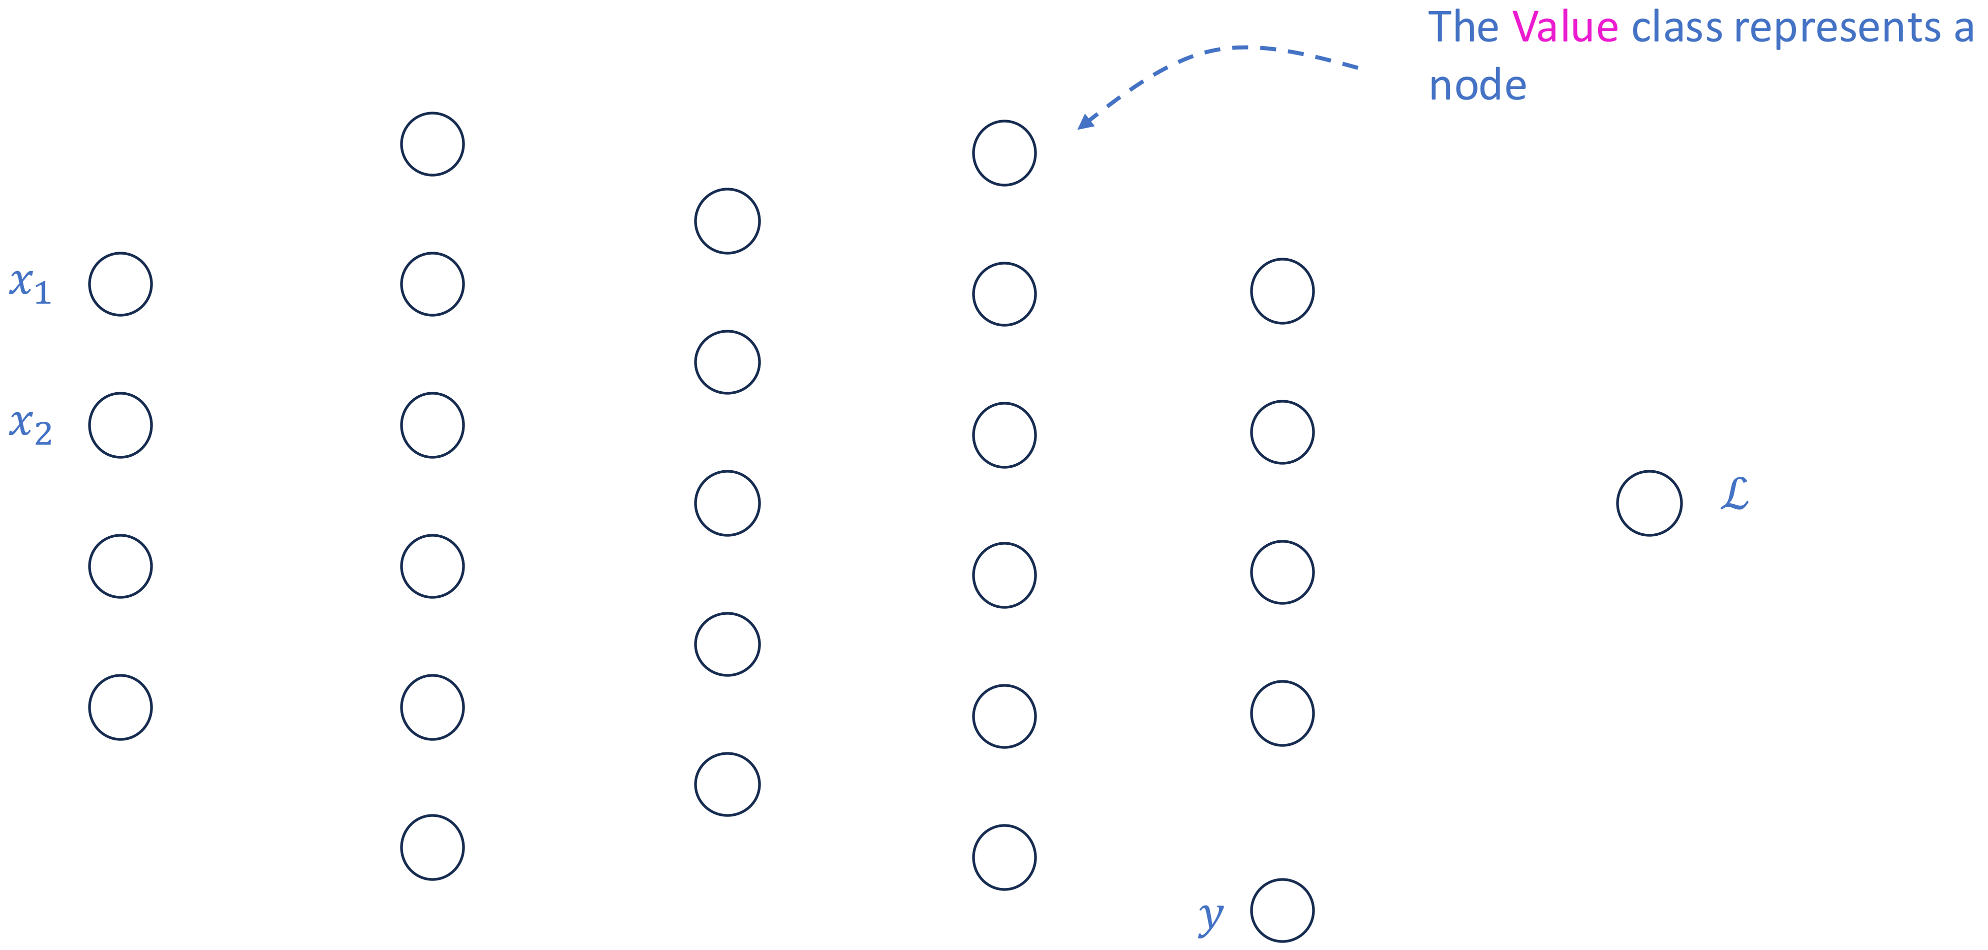

Let's build it by hand!

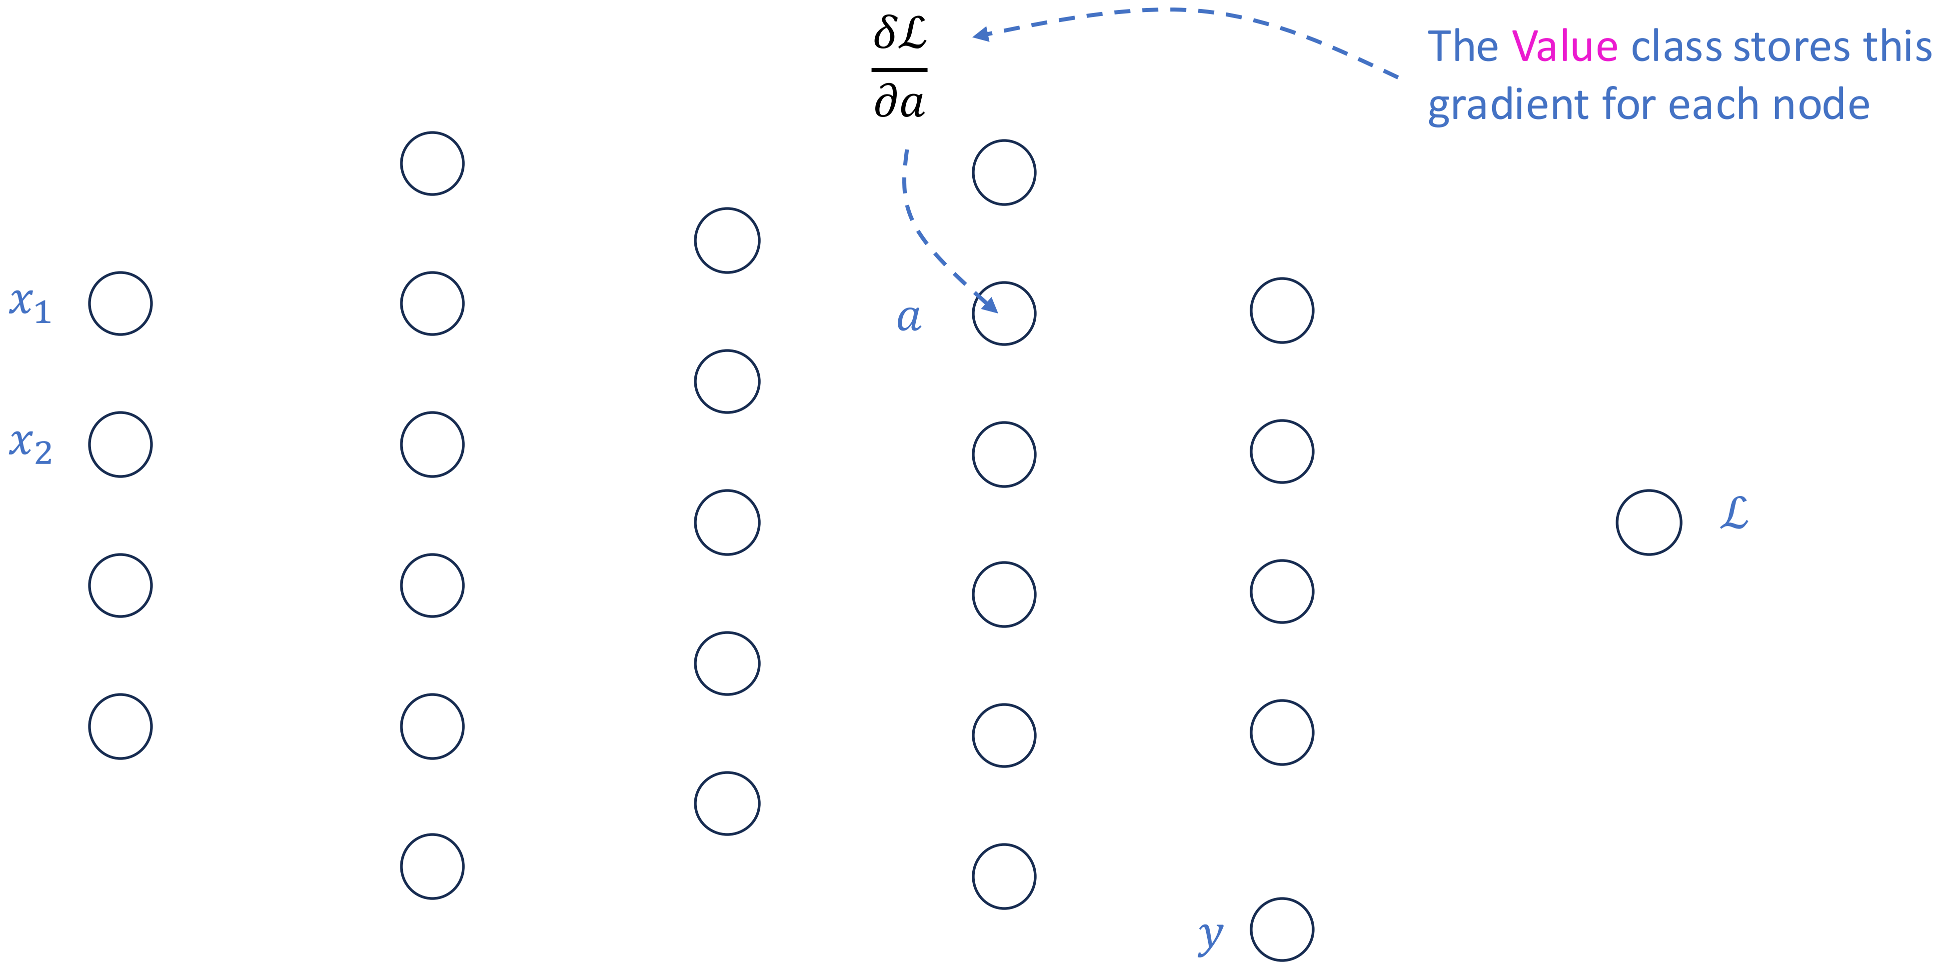

Let’s build it by hand!

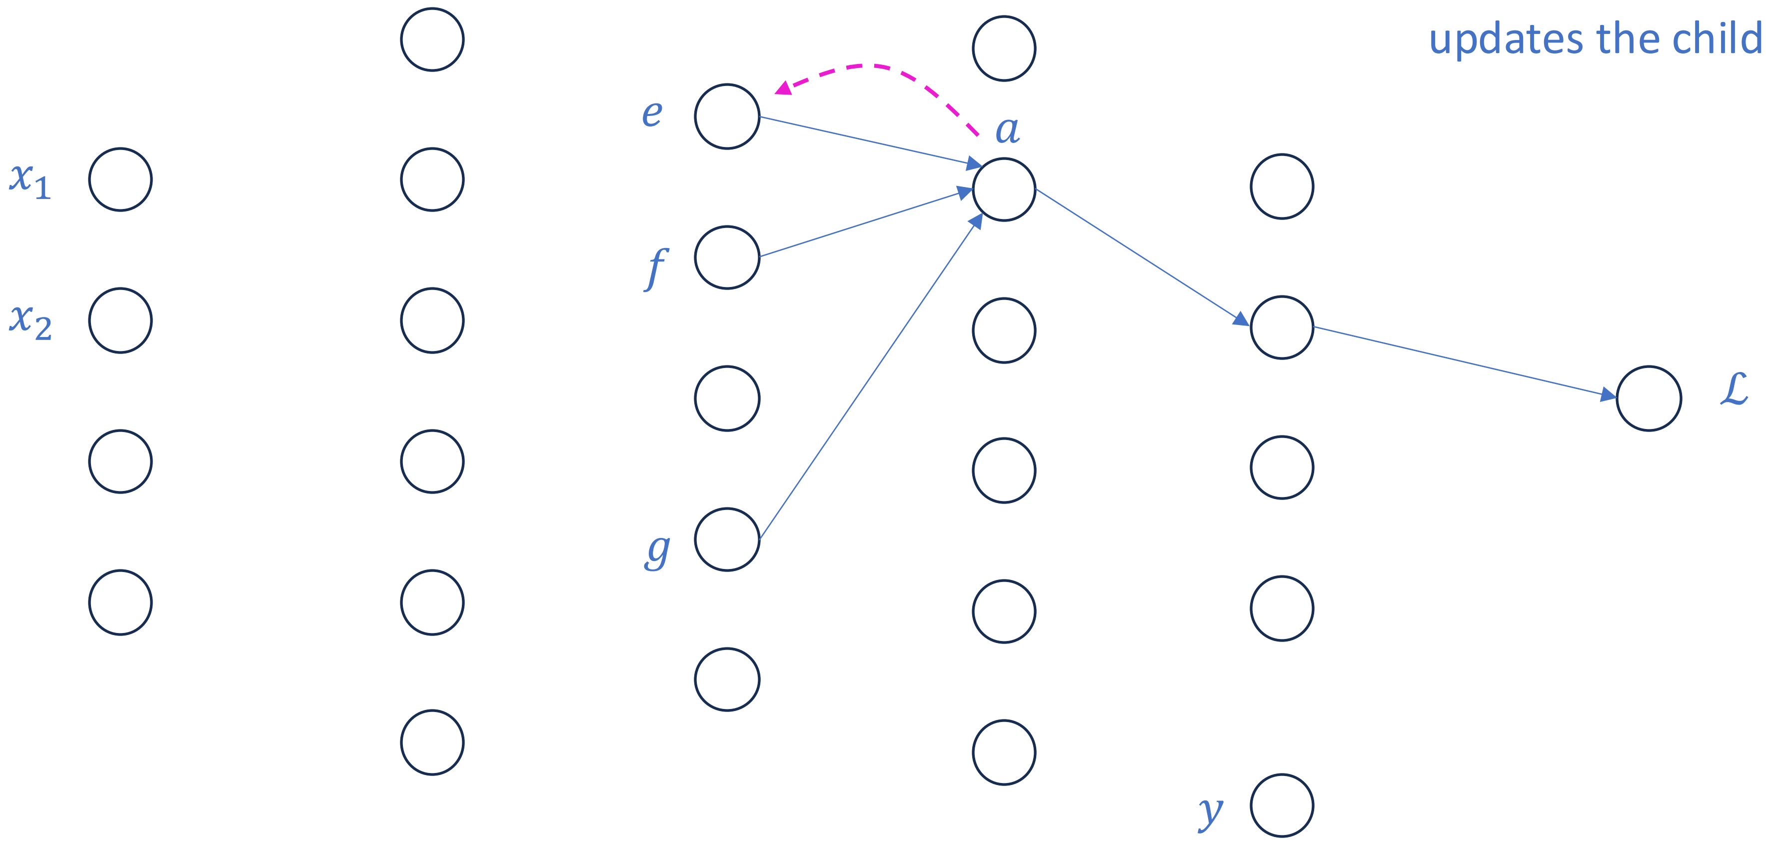

The Value class contains a “backward” function that updates the children’s gradient

Let's build it by hand!

e.grad += $\frac{\partial a}{\partial e}$ * a.grad

$$\frac{\partial \mathcal{L}}{\partial e} = \frac{\partial \mathcal{L}}{\partial a}\frac{\partial a}{\partial e} + \frac{\partial \mathcal{L}}{\partial b}\frac{\partial b}{\partial e} + \cdots$$

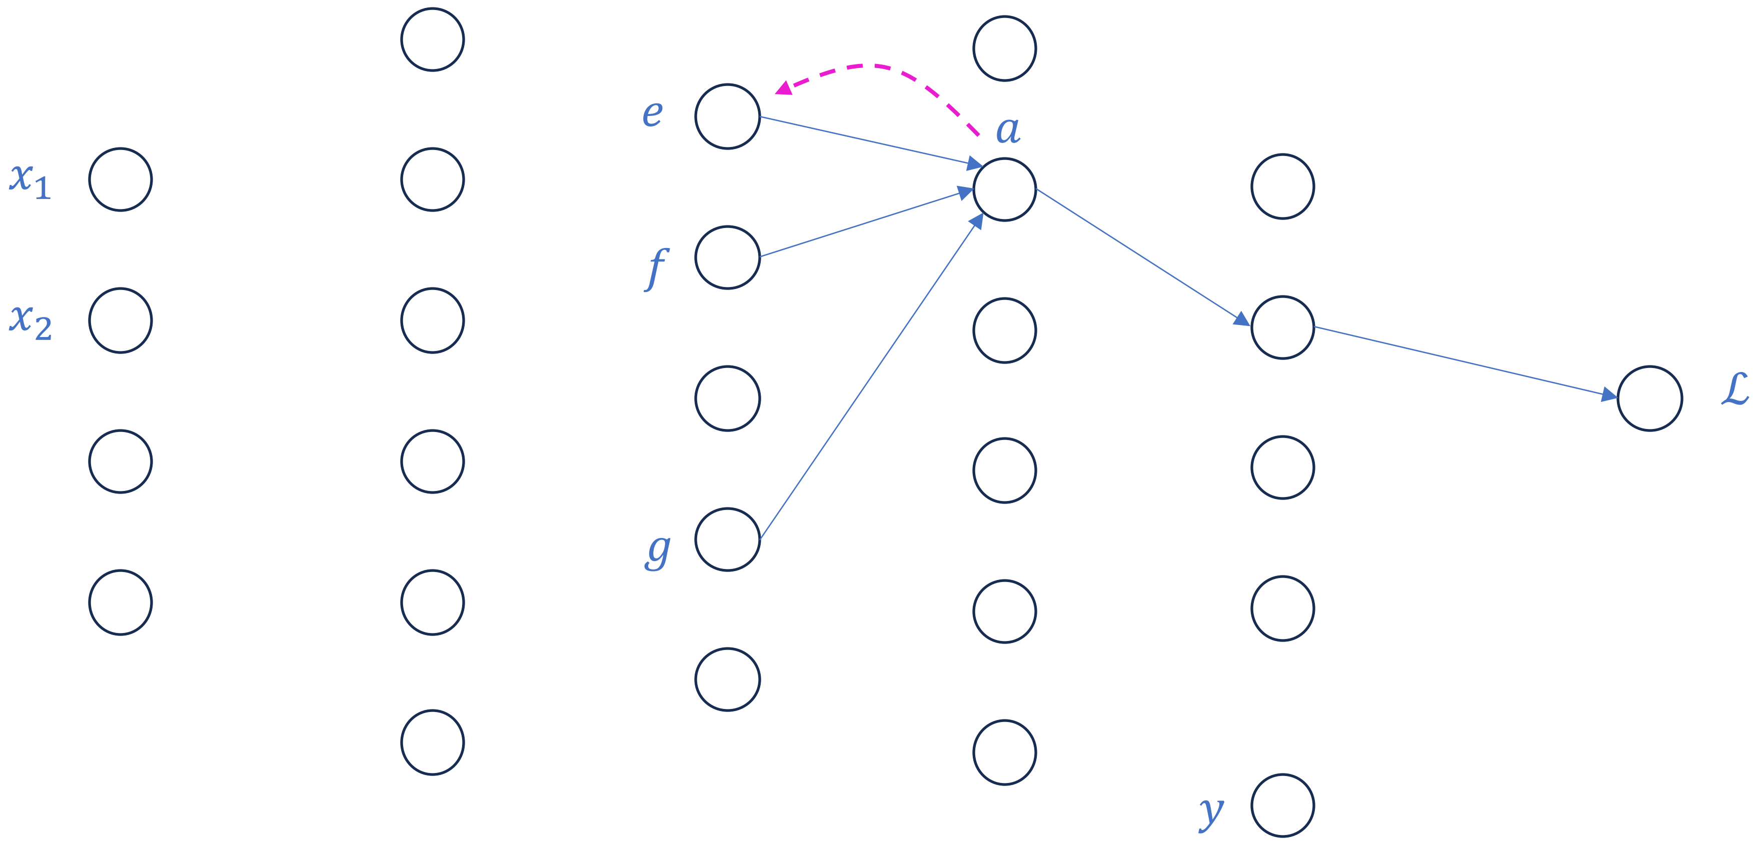

Let's build it by hand!

$$\texttt{e.grad += } \frac{\partial a}{\partial e} \texttt{ * a.grad}$$

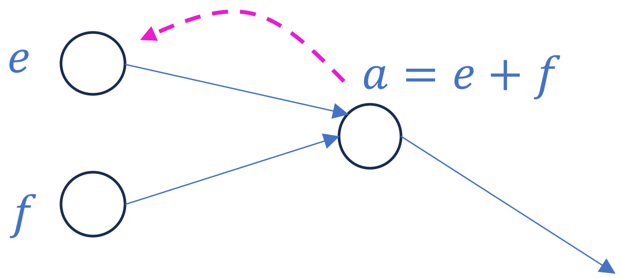


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0


    def __add__ (self, other):
        out = Value(self.data + other.data)
        return out
```

Let's build it by hand!

$$\texttt{e.grad +=} \frac{\partial a}{\partial e} \texttt{ * a.grad}$$

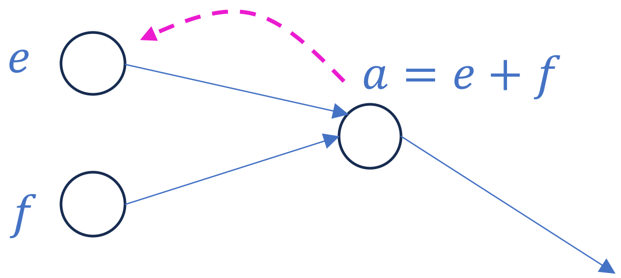


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __add__ (self, other):
        out = Value(self.data + other.data)
        def _backward():
            self.grad += ???
            other.grad += ???
        out.backward = _backward
        return out
```

Let's build it by hand!

$$\texttt{e.grad += } \frac{\partial a}{\partial e} \texttt{ * a.grad}$$

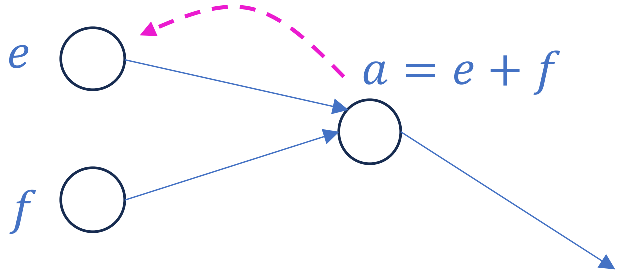


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __add__ (self, other):
        out = Value(self.data + other.data)
        def _backward():
            self.grad += 1.0 * out.grad
            other.grad += 1.0 * out.grad
        out.backward = _backward
        return out
```

Let’s build it by hand!

e.grad += $\frac{\partial a}{\partial e}$ * a.grad

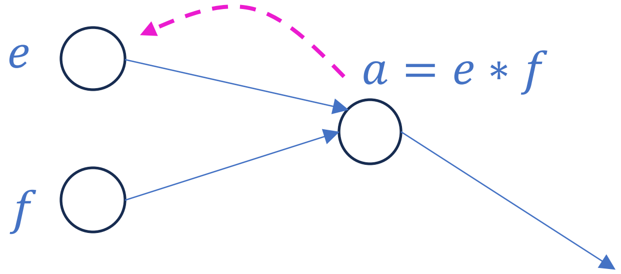


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __mul__ (self, other):
        out = Value(self.data * other.data)
        def _backward():
            self.grad += ??? * out.grad
            other.grad += ??? * out.grad
        out.backward = _backward
        return out
```

Let’s build it by hand!

e.grad += $\frac{\partial a}{\partial e}$ * a.grad

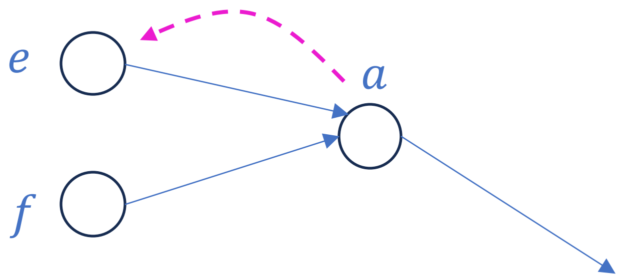


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __mul__ (self, other):
        out = Value(self.data * other.data)
        def _backward():
            self.grad += other.data * out.grad
            other.grad += self.data * out.grad
        out.backward = _backward
        return out
```

$$a = e * f$$

$$\frac{\partial a}{\partial e} = f$$

Let's build it by hand!

e.grad += $\frac{\partial a}{\partial e}$ * a.grad

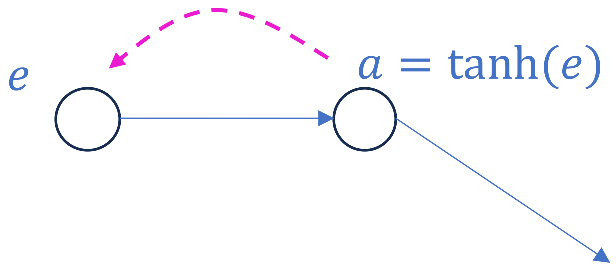


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __tanh__ (self):
        out = Value(???)
        def _backward():
            self.grad += ??? * out.grad
        out.backward = _backward
        return out
```

$$a = \tanh(e)$$

$$\frac{\partial a}{\partial e} = ???$$

Let’s build it by hand!

$$\texttt{e.grad += } \frac{\partial a}{\partial e} \texttt{ * a.grad}$$

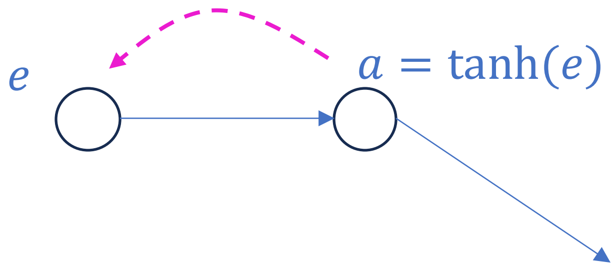


```
class Value:
    def __init__(self, data):
        self.data = data
        self.grad = 0.0
        self.backward = lambda : None

    def __tanh__ (self):
       t = (math.exp(2*x) - 1)/(math.exp(2*x) + 1)
        out = Value(t, (self, ), 'tanh')
        def _backward():
            self.grad += (1.0 - t**2) * out.grad
        out.backward = _backward
        return out
```

$$a = \tanh(e)$$

$$\frac{\partial a}{\partial e} = 1 - a^2$$

The lol() function

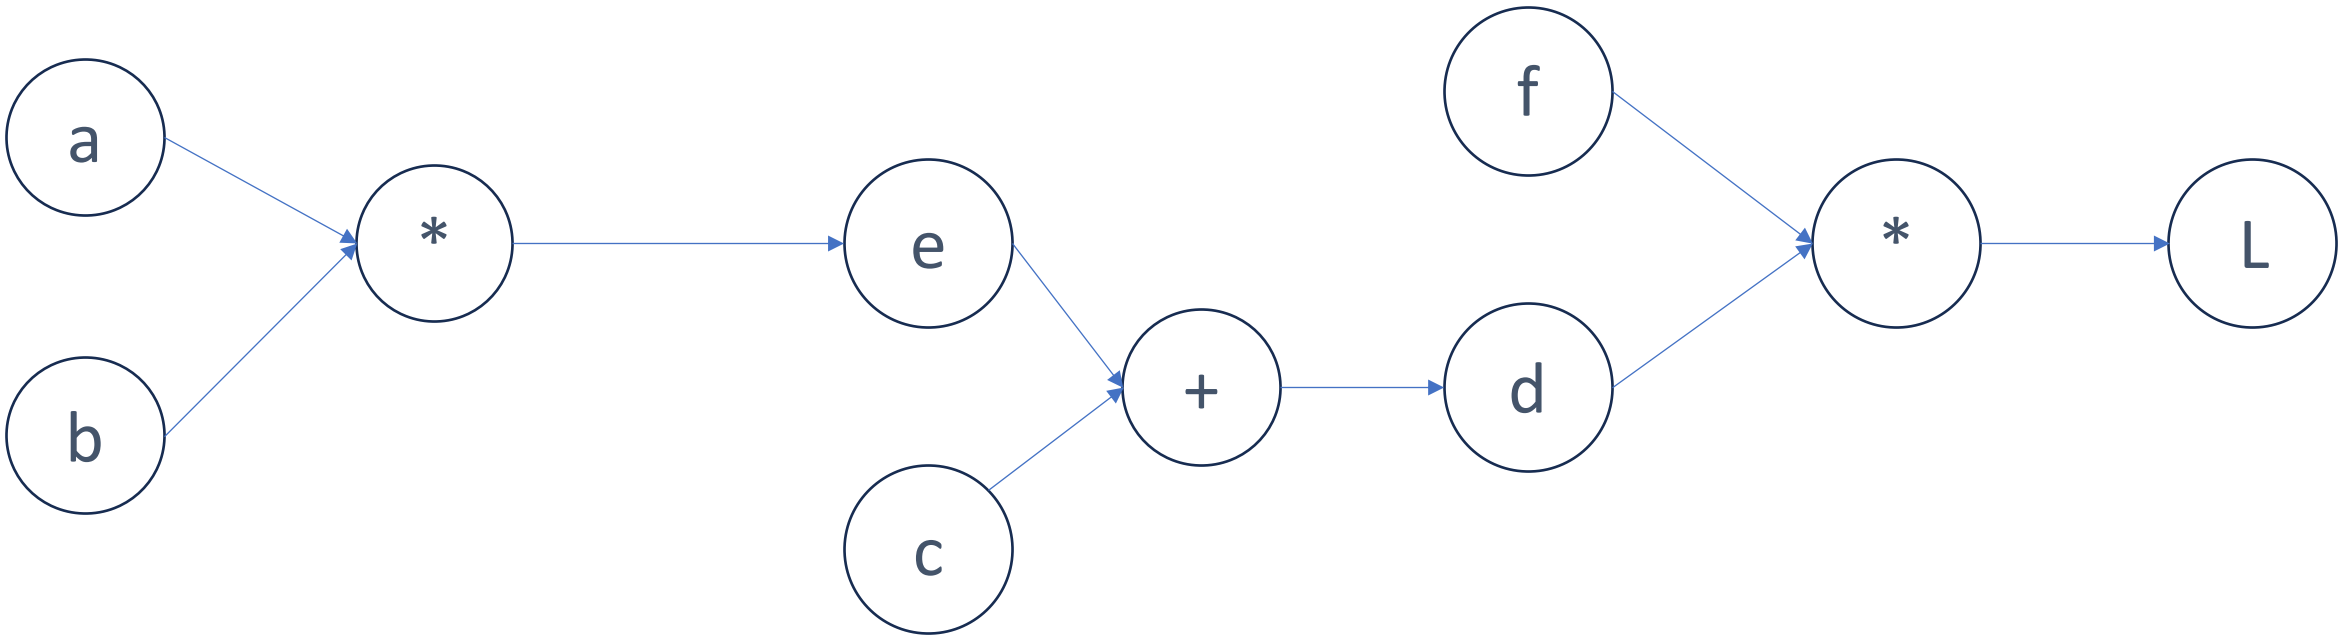

The lol() function

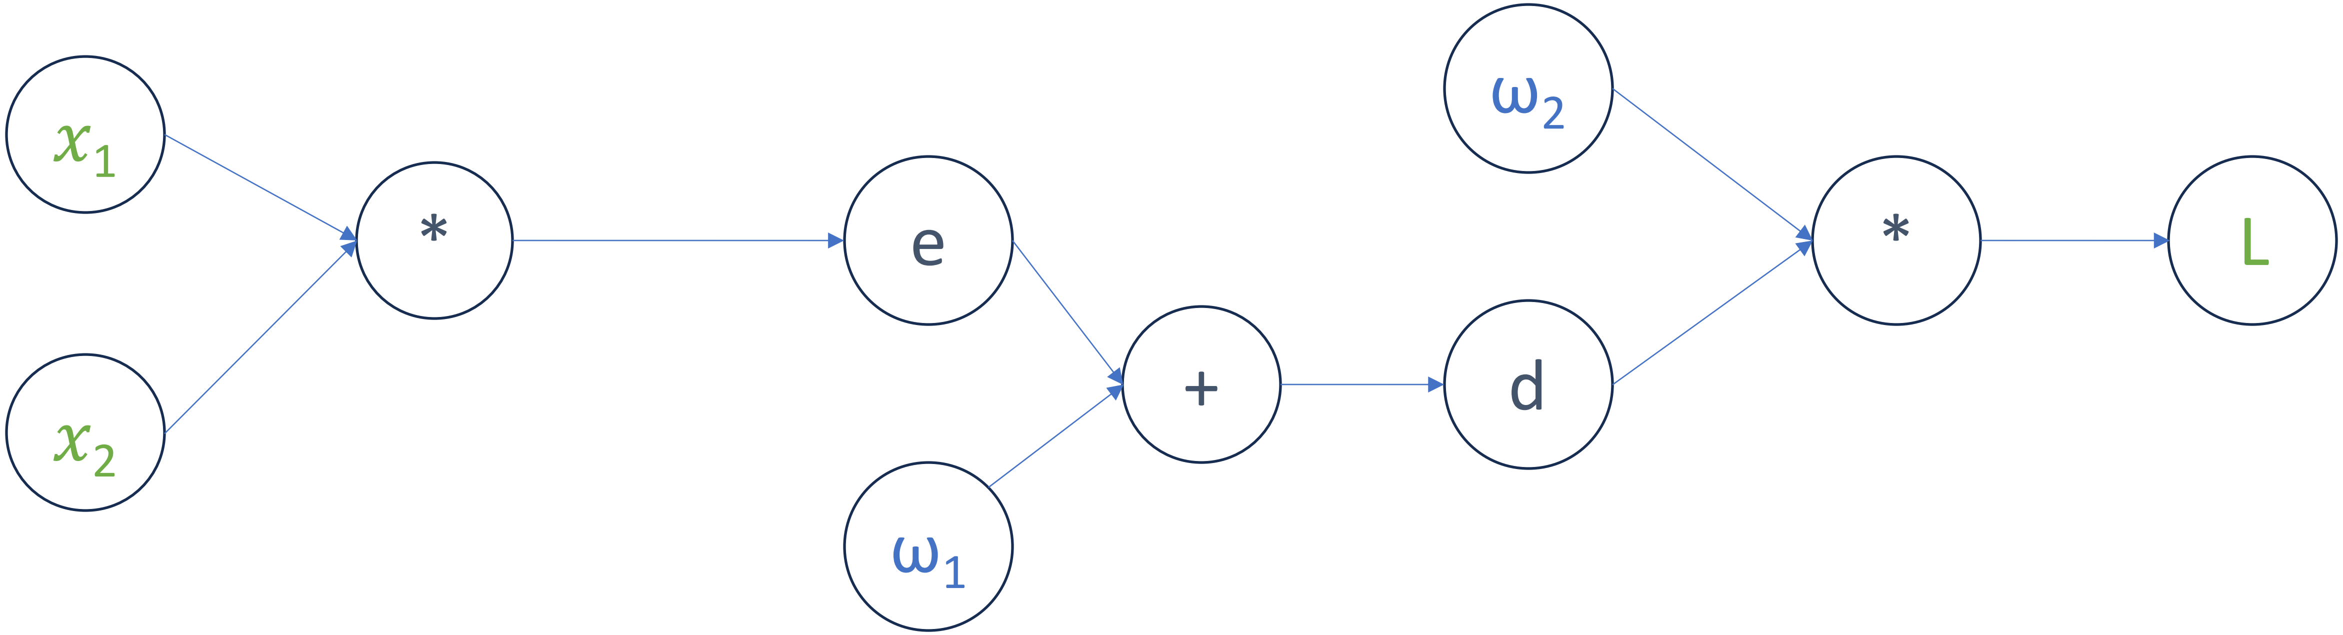

The lol() function

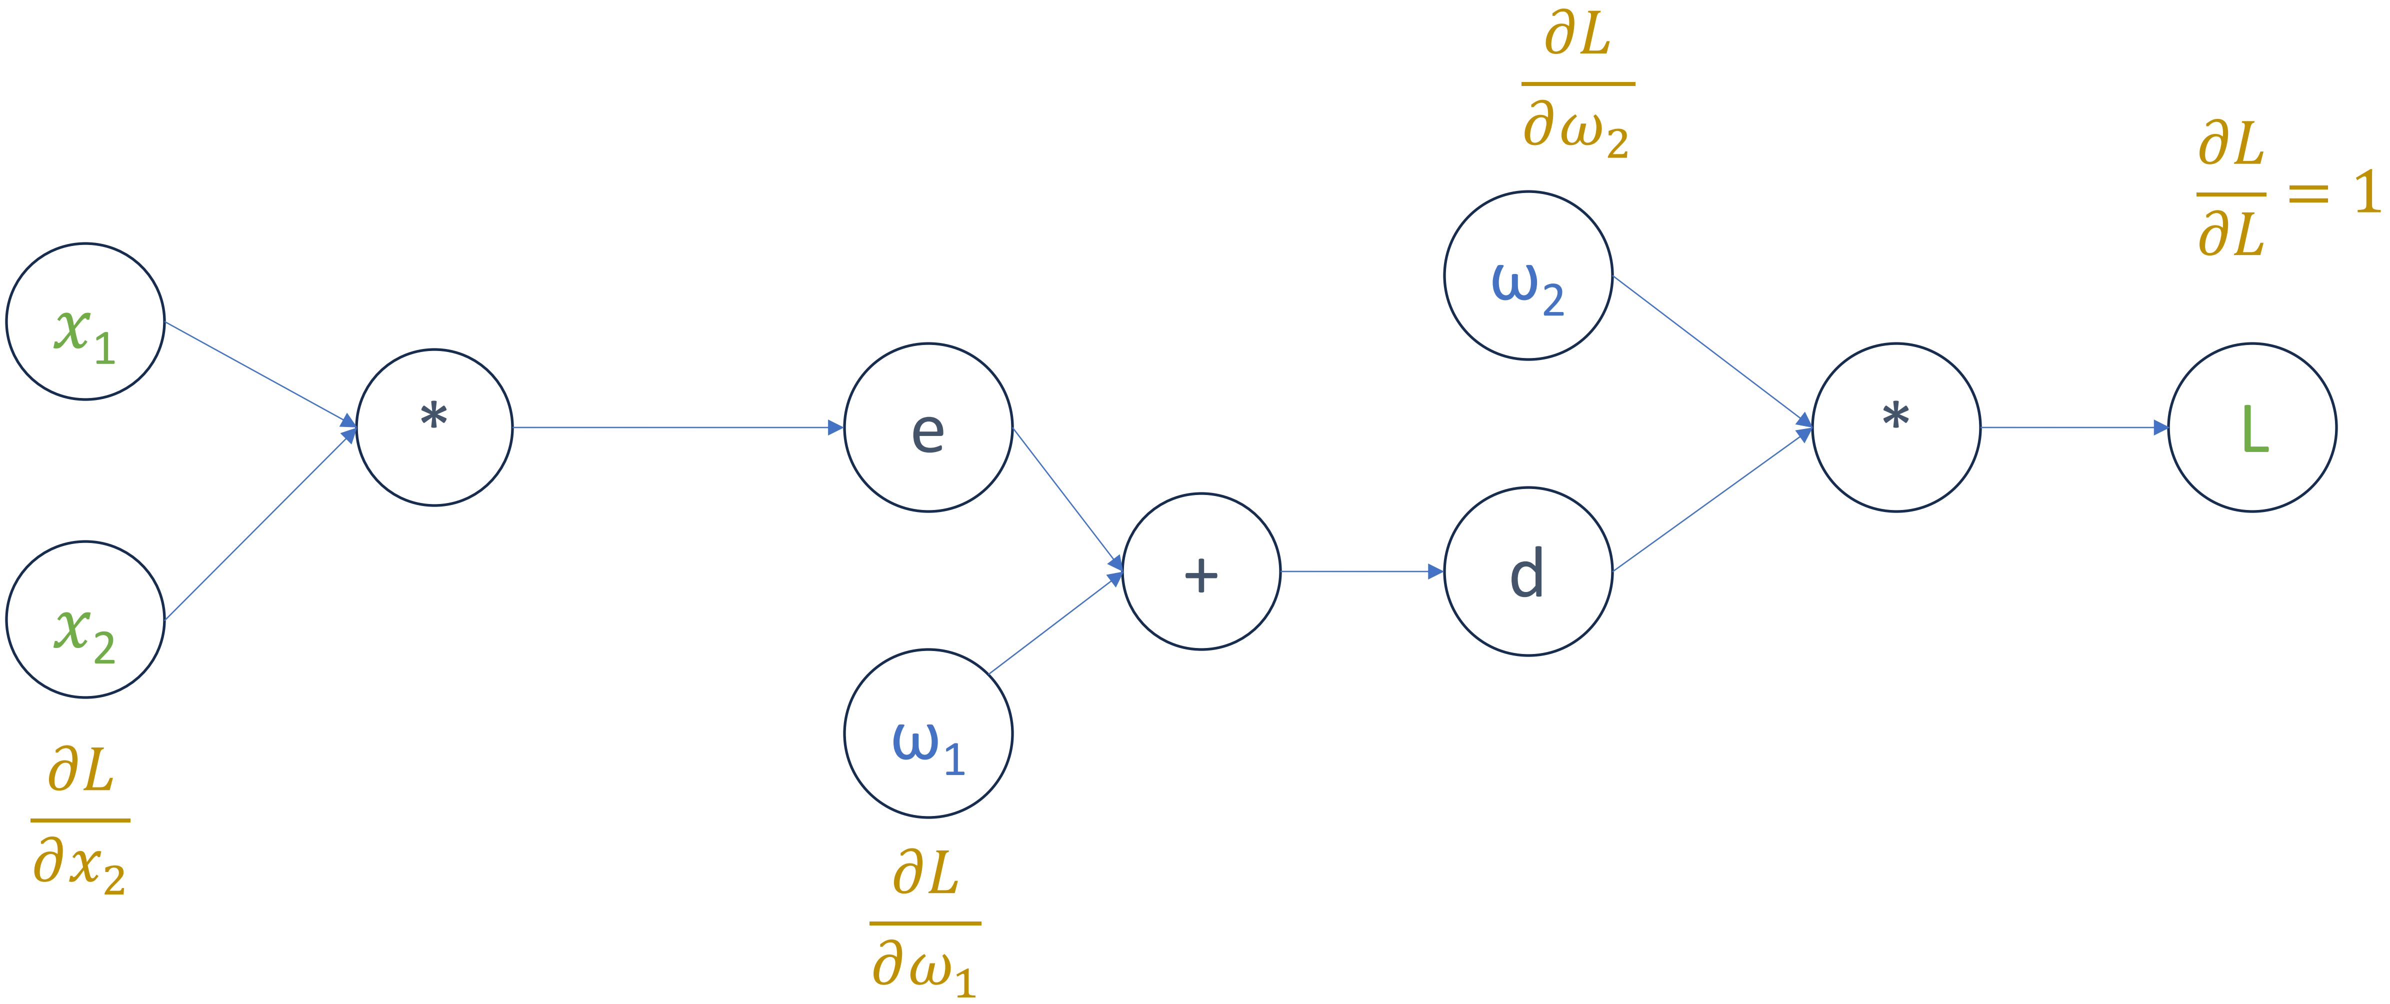

Gradient descent for the lol() function

$$\omega_1 = \omega_1 - \eta \frac{\partial L}{\partial \omega_1}$$

$$\omega_2 = \omega_2 - \eta \frac{\partial L}{\partial \omega_2}$$

The lol() function

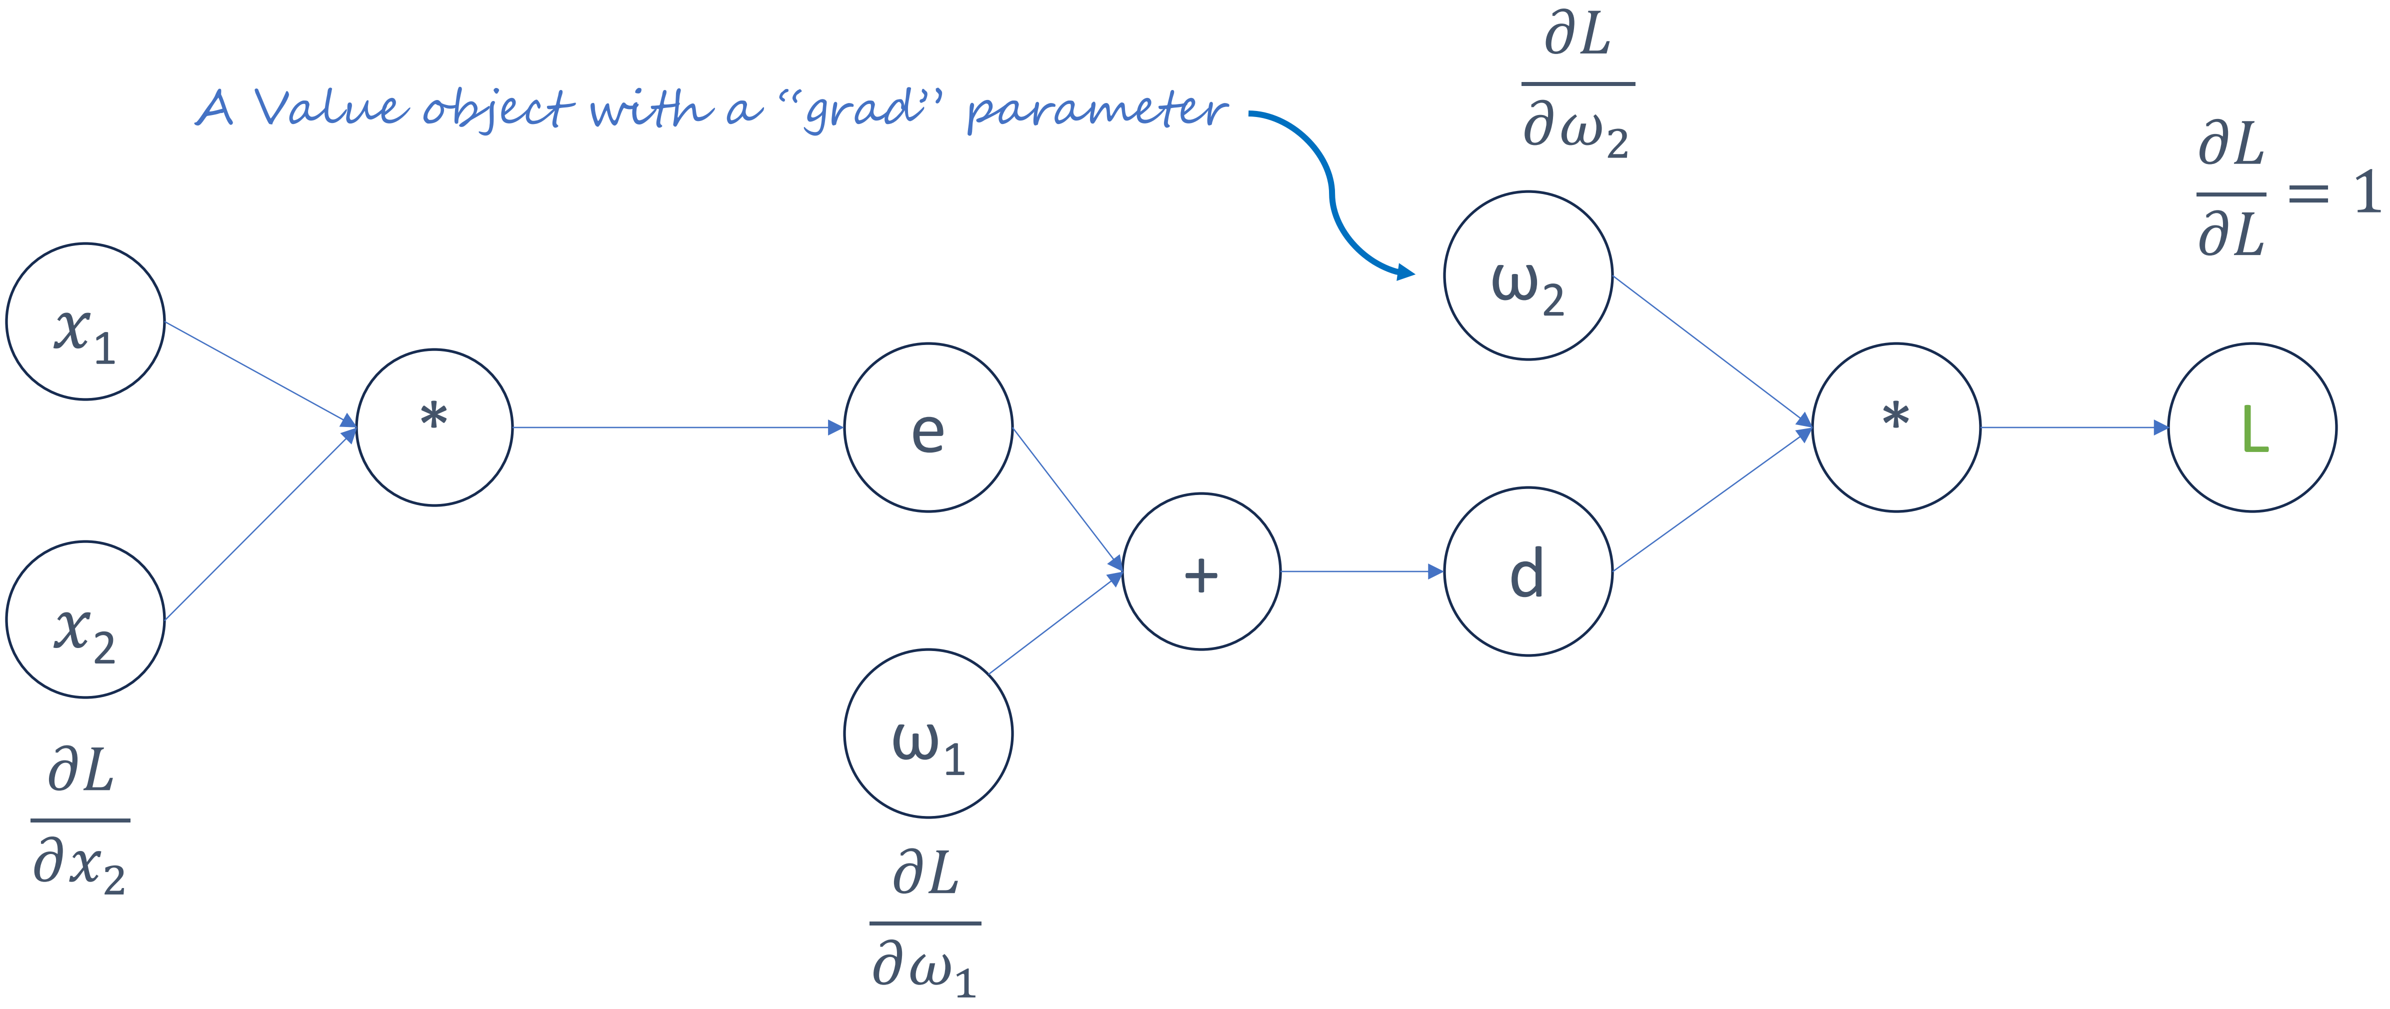

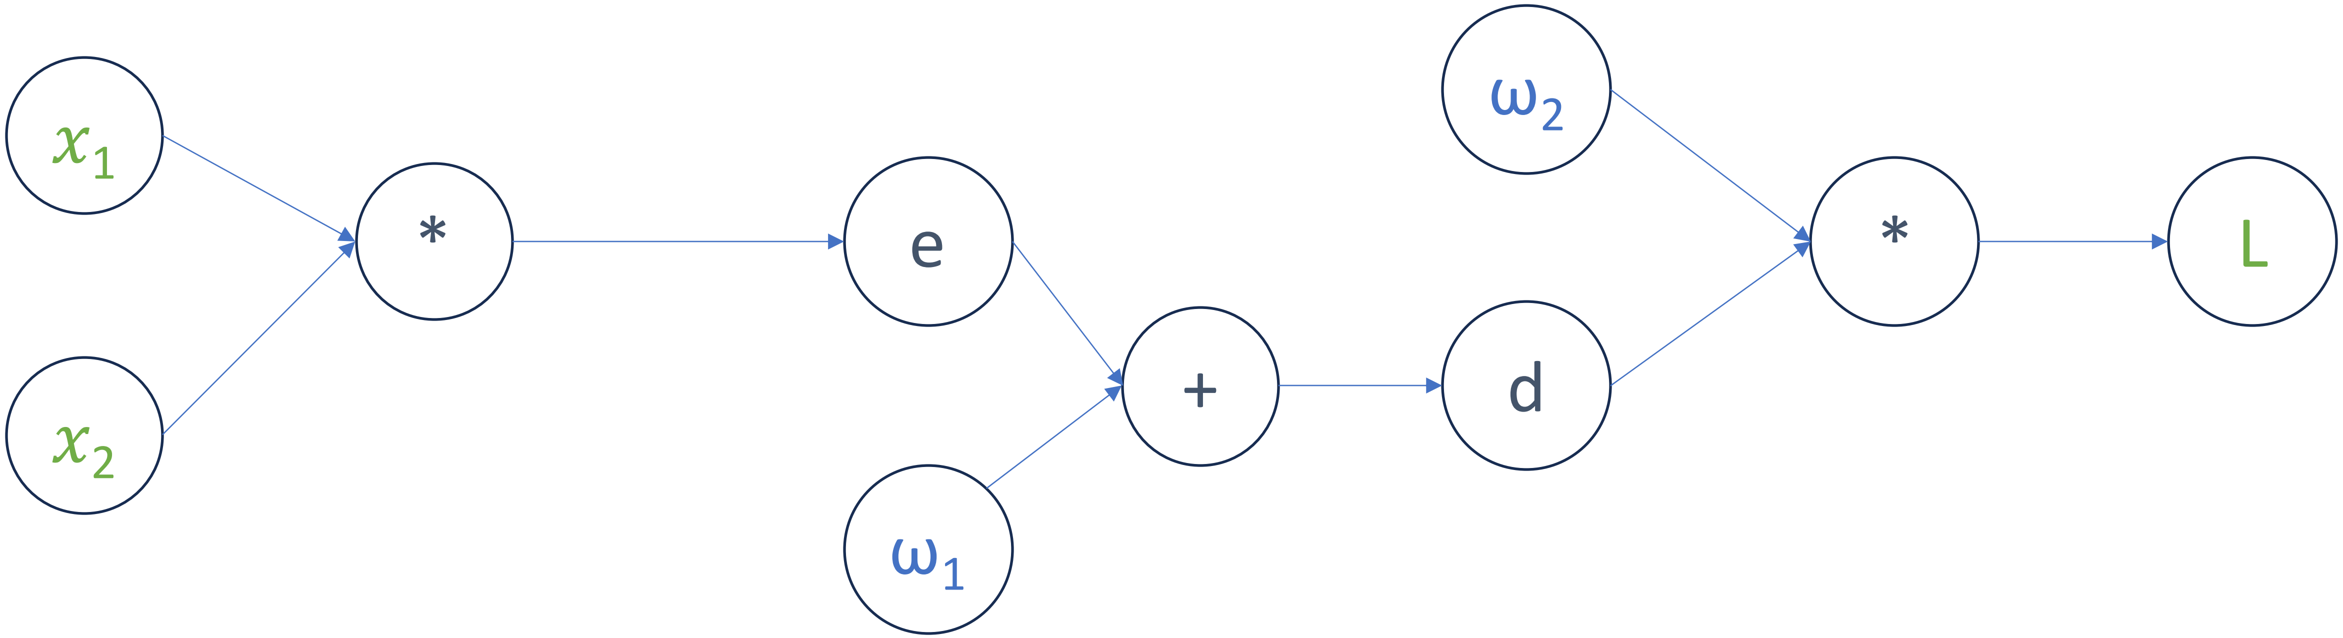


```python
class Neuron:

  def __init__(self, nin):
    self.w = [Value(random.uniform(-1,1)) for _ in range(nin)]
    self.b = Value(random.uniform(-1,1))

  def __call__(self, x):
    # w * x + b
    act = ???
    out = act.tanh()
    return out

  def parameters(self):
    return self.w + [self.b]
```

Mean square error loss

$$\ell = \sum (y - \hat{y})^2$$

```
n = MLP(3, [4, 4, 1])

ypred = [n(x) for x in xs]

loss = sum([(a-b)**2 for (a,b) in zip(ypred, ys)])
```

Compute graphs become huge very quickly

This is a 3-layer MLP with 41 parameters!

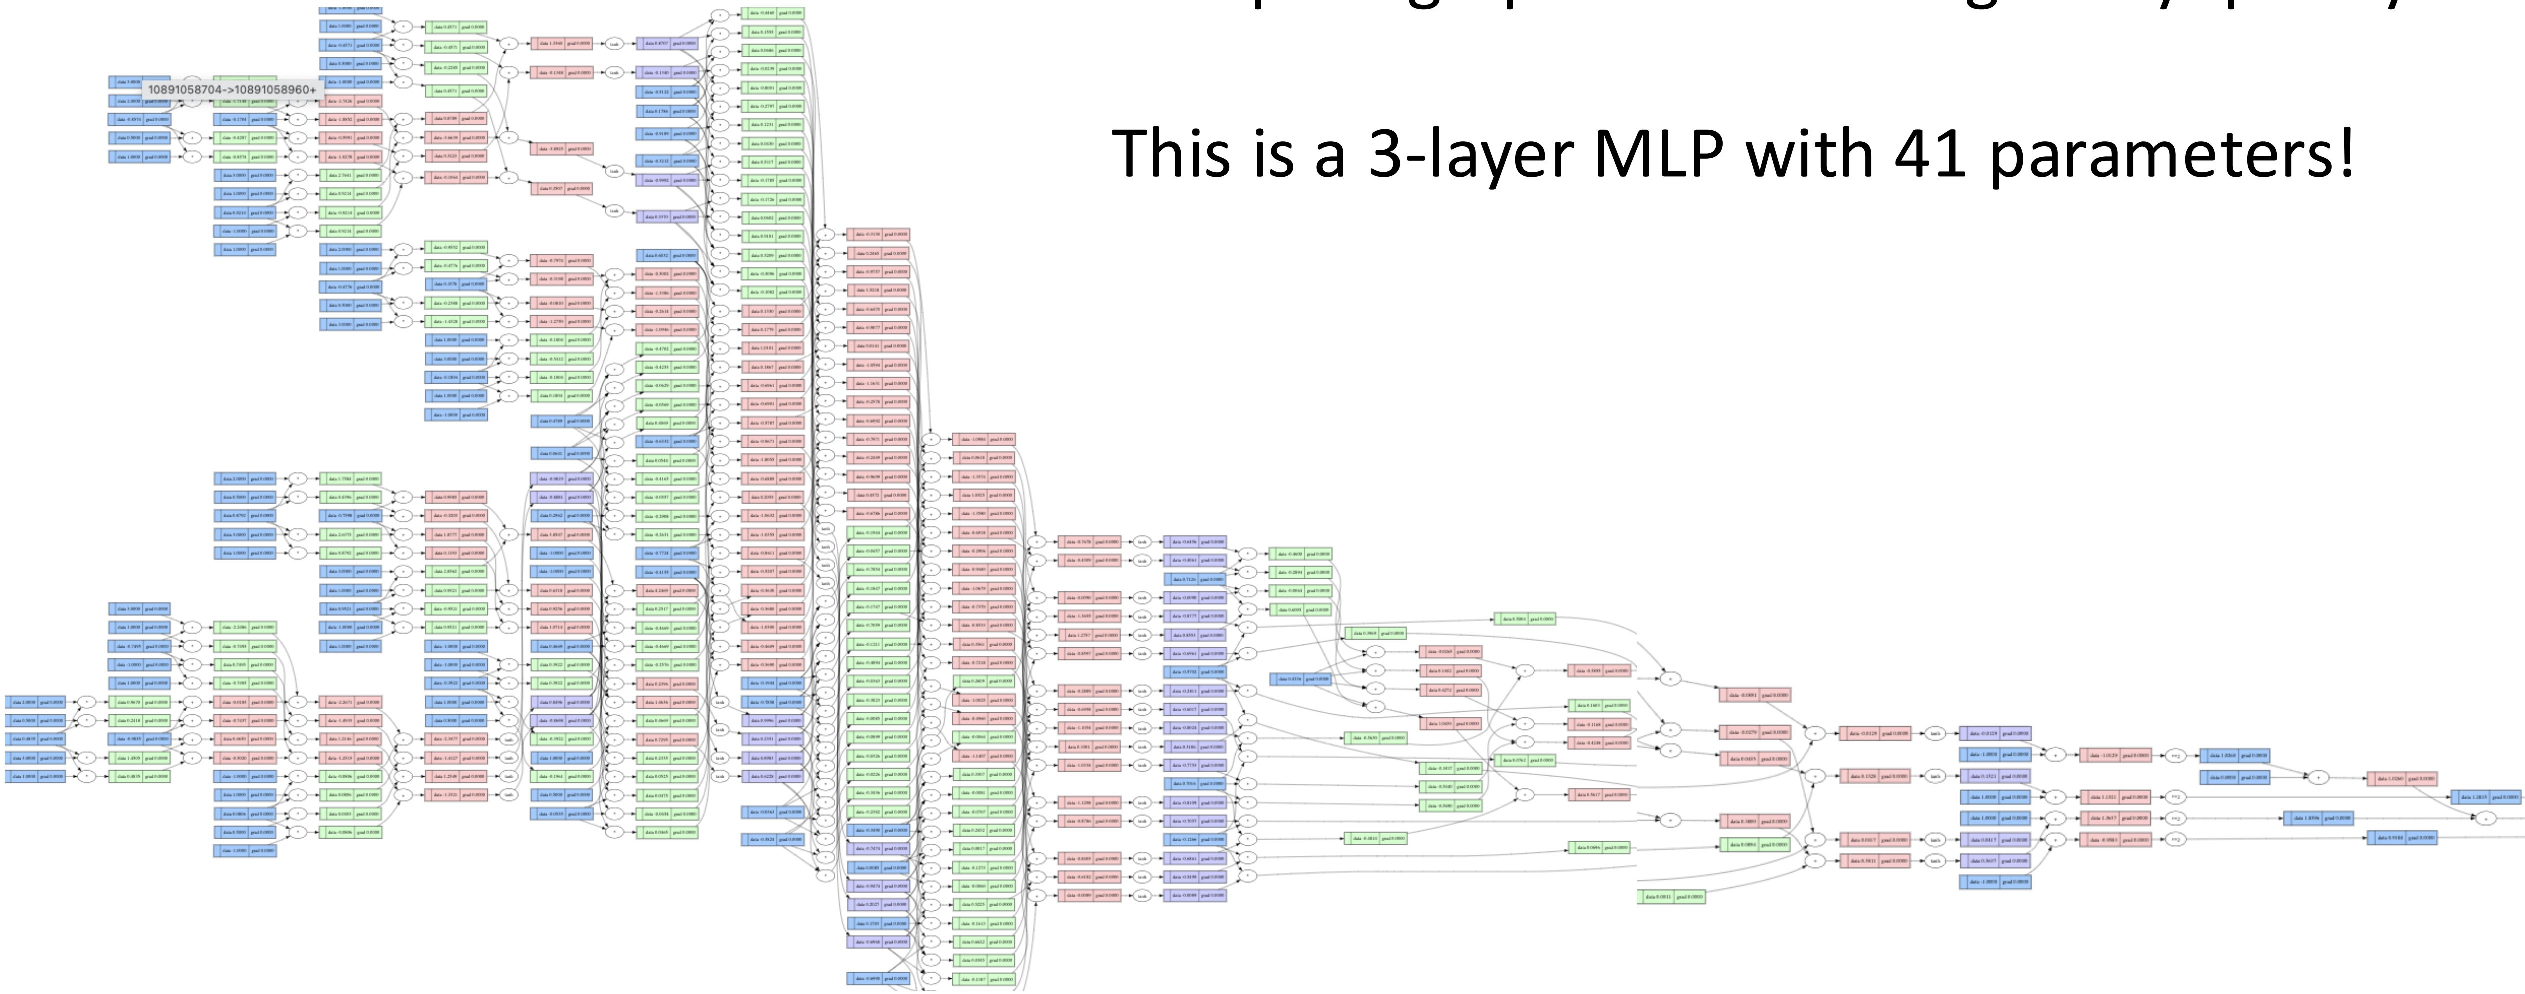

The forward pass is critical to update the values in the network

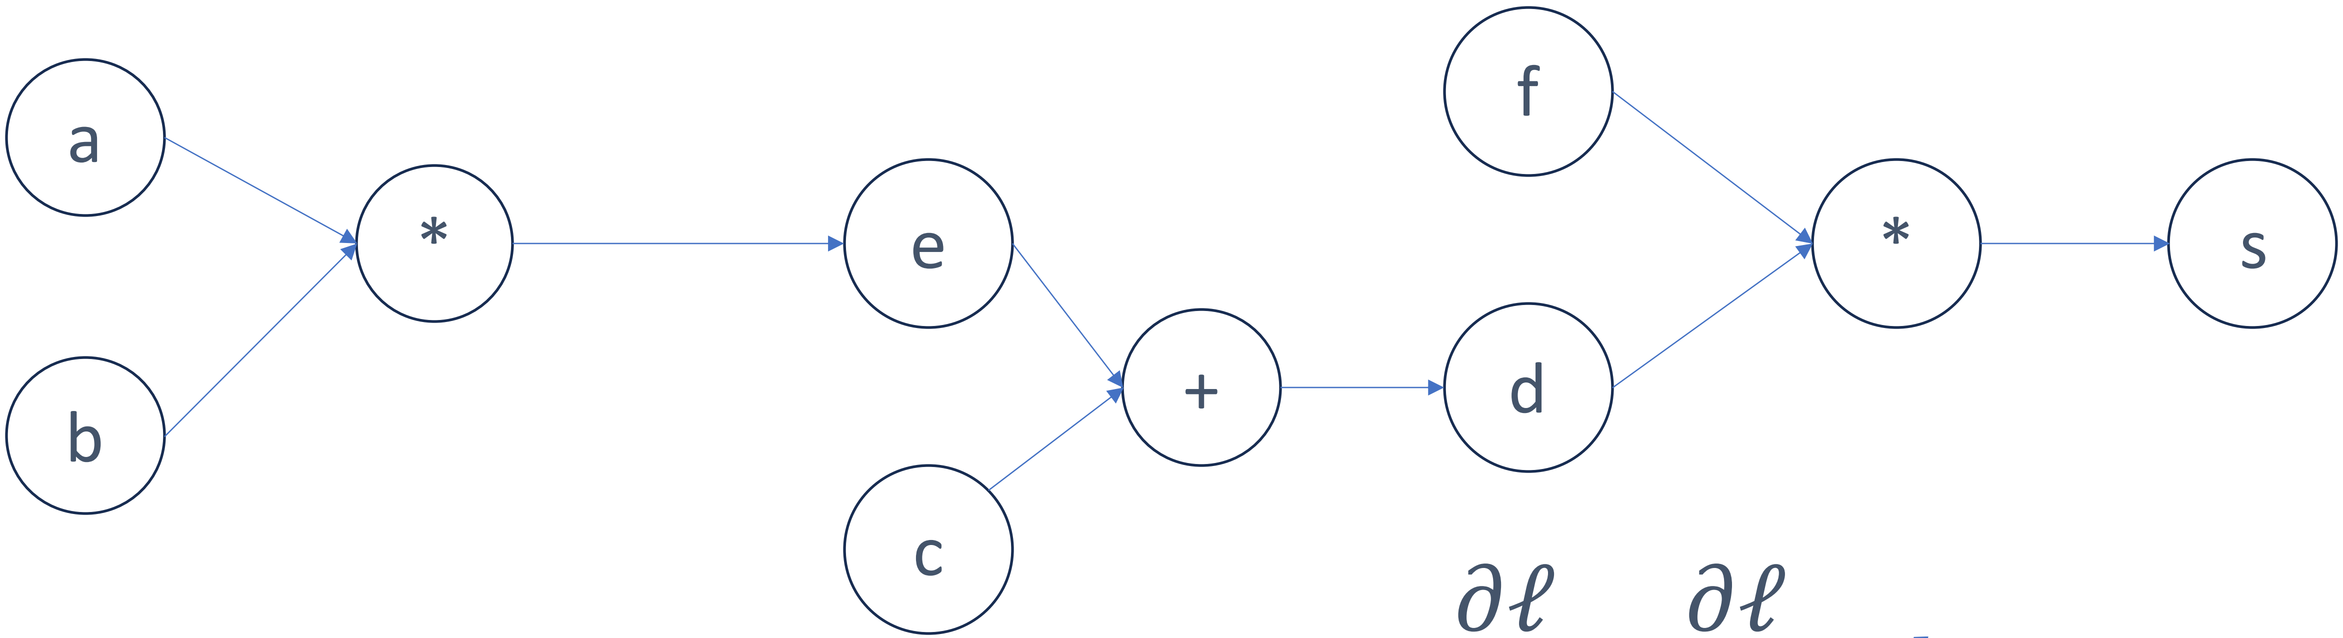

$$\frac{\partial \ell}{\partial f} = \frac{\partial \ell}{\partial s} * d$$

Stochastic gradient descent

Iterate over batches of samples instead of the whole dataset at once

1. Scales to large datasets (that don’t fit in memory)
2. Adds randomness to the process
3. Adds implicit regularization

Conclusion

Neural networks are compute graphs.

Gradient descent minimizes a loss function over the network’s parameters.

Back-propagation allows efficient learning (tuning of the network).

Let’s practice with a simple Multi-Layer Perception (MLP).

Universal approximation theorem

For any continuous function, there exists a shallow network that can approximate this function to any specified precision.

Universal approximation theorem

Universal approximation theorem (Uniform non-affine activation, arbitrary depth, constrained width). Let $\mathcal{X}$ be a compact subset of $\mathbb{R}^d$. Let $\sigma : \mathbb{R} \to \mathbb{R}$ be any non-affine continuous function which is continuously differentiable at at least one point, with nonzero derivative at that point. Let $\mathcal{N}^{\sigma}_{d,D:d+D+2}$ denote the space of feed-forward neural networks with d input neurons, D output neurons, and an arbitrary number of hidden layers each with $d + D + 2$ neurons, such that every hidden neuron has activation function σ and every output neuron has the identity as its activation function, with input layer ϕ and output layer ρ. Then given any $\varepsilon > 0$ and any $f \in C(\mathcal{X}, \mathbb{R}^D)$, there exists $\hat{f} \in \mathcal{N}^{\sigma}_{d,D:d+D+2}$ such that

$$\sup_{x \in \mathcal{X}} \left\| \hat{f}(x) - f(x) \right\| < \varepsilon.$$

In other words, $\mathcal{N}$ is dense in $C(\mathcal{X}; \mathbb{R}^D)$ with respect to the topology of uniform convergence.

No free-lunch theorem

Every learning algorithm is as good as any other when averaged over all sets of problems.

You can’t just learn « purely from data » without bias.

Wolpert, D. H.; Macready, W. G. (1997). "No Free Lunch Theorems for Optimization". *IEEE Transactions on Evolutionary Computation*. **1**: 67–82.